

# Capstone Project - Udacity Machine Learning Engineer Nanodegree

## Definition

### Project Overview

Deep learning is the new big trend in machine learning. It had many recent successes in computer vision, automatic speech recognition and natural language processing. The goal of this project is to give you a hands-on introduction to deep learning. To do this, we will build a Cat/Dog image classifier using a deep learning algorithm called convolutional neural network (CNN) and a ImageNet dataset.

Deep Learning attempts to use large volumes of unstructured data, such as images and audio clips, to learn hierarchical models which capture the complex structure in the data and then use these models to predict properties of previously unseen data. For example, DL has proven extremely successful at learning hierarchical models of the visual features and concepts represented in handheld camera images and then using those models to automatically label previously unseen images with the objects present in them.

Humans are very good at recognizing visual patterns--including the faces of people you know--while it's maddeningly complex for a computer to "understand" something visually, dot-by-dot. We will build a Cat/Dog image classifier inspired from Kaggle competition (<https://www.kaggle.com/c/dogs-vs-cats>).



## Problem Statement

Our goal is to build a machine learning algorithm capable of detecting the correct animal (cat or dog) in new unseen images.

The objectives of this project are to complete the following tasks with Deep Learning:

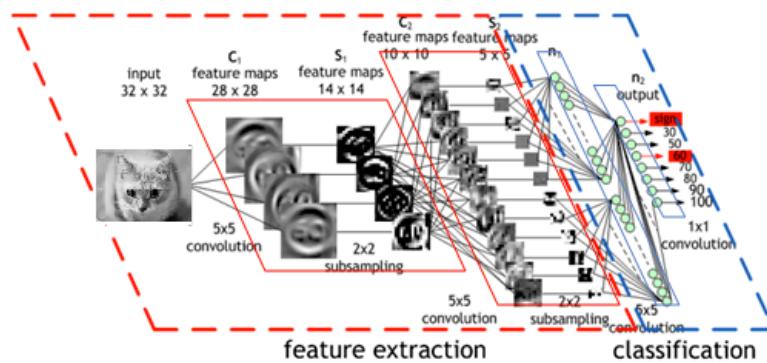
1. Download and preprocess dog/cat images from ImageNet and move them to a database.
2. Build and train a convolutional neural network (CNN) using Caffe for classifying images. Here is an intuitive explanation of the [CNN \(<https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>\)](https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/).
3. Evaluate the classification performance of a trained CNN under different training parameter configurations.
4. Modify the network configuration to improve classification performance.
5. Visualize the features that a trained network has learned.

The intended solution for this problem is to classify new unseen images with highest possible accuracy using a convolutional neural network completely trained in the project from scratch.

## Metrics

In deep learning, the final layer of a neural network used for classification can often be interpreted as a logistic regression. In this context, one can see a deep learning algorithm as multiple feature learning stages, which then pass their features into a logistic regression that classifies an input.

In the next sections of this project, we will look carefully at models we create. We will define two closely related networks in one file: the network used for training and the network used for testing. Both networks output the softmax\_loss layer, which in training is used to compute the loss function and to initialize the backpropagation, while in validation this loss is simply reported. The testing network also has a second output layer, accuracy, which is used to report the accuracy on the test set.



Because each training example belongs to one class, we assign the training example to the class whose probability is the highest and generate a vector of binary values. Then, we compare the class assigned by the classifier and its true class to evaluate the accuracy. The official accuracy metric used:

$$acc = \frac{NumCorrectPredictions}{numExamples}$$

This process is repeated many thousands of times until the network converges to a stable average classification accuracy across all the training images.

## Analysis

## Data Exploration

For this project, a subset of the [ImageNet \(http://www.image-net.org/\)](http://www.image-net.org/) dataset has been used. ImageNet, a largescale ontology of images built upon the backbone of the WordNet structure. ImageNet aims to populate the majority of the 80,000 synsets of WordNet with an average of 500-1000 clean and full resolution images. This will result in tens of millions of annotated images organized by the semantic hierarchy of WordNet.

The images used in this project are from two different categories, cats and dogs, downloaded from Imagenet. The training data and validation data are downloaded from ImageNet and saved as jpg files, and they are stored on the disk drive like:

- /home/ubuntu/Data/dog\_cat\_32/train
- /home/ubuntu/Data/dog\_cat\_32/test
- /home/ubuntu/Data/dog\_cat\_32/val

To minimize the training time, all of the images have been resized to 32x32 pixels. The training and validation input are described in train.txt and val.txt as text listing all the files and their labels. These files simply list the relative filename of each image tab separated from a natural number representing the class the image belongs to. For example, train.txt contains the following rows:

```
cat/cat_0_32.jpg 0 cat/cat_1000_32.jpg 0 dog/dog_9990_32.jpg 1 dog/dog_9991_32.jpg 1 ...
```

## Exploratory Visualization

Below are sample images from both categories. The cats category includes domestic cats as well as large breeds like lions and tigers. The dog category is comprised of domestic dogs including pugs, basenji and great pyrenees. There are approximately 13,000 images in total.

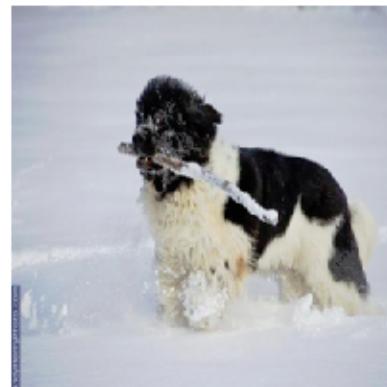


	Image name	Label
--	------------	-------

## Visualization of class labels

```
In [14]: import pandas as pd  
  
%matplotlib inline  
  
user_cols = ['Image name','Label']  
data_dir = "/home/ubuntu/Data/dog_cat_32/train/train.txt"  
  
data = pd.read_csv(data_dir, sep = ' ', names = user_cols, header = None)  
print "Data read successfully!"
```

Data read successfully!

```
In [15]: data.shape
```

Out[15]: (12904, 2)

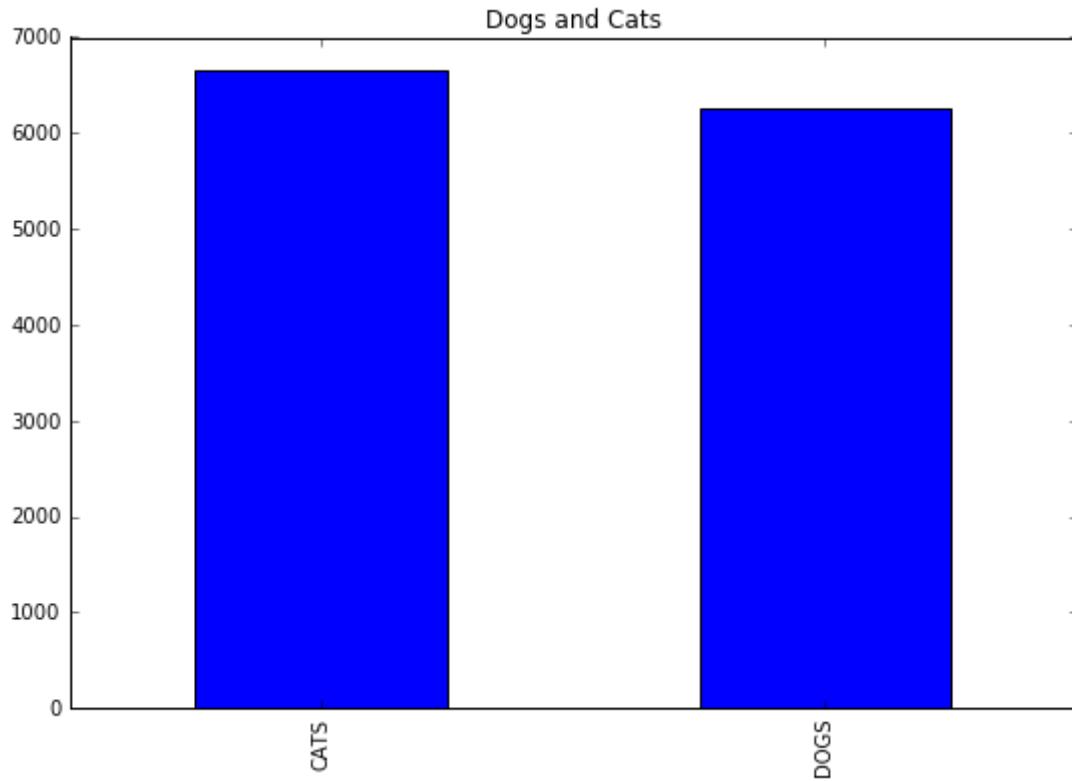
```
In [16]: data.head()
```

Out[16]:

	Image name	Label
0	cat/cat_0_32.jpg	0
1	cat/cat_1000_32.jpg	0
2	cat/cat_1001_32.jpg	0
3	cat/cat_1002_32.jpg	0
4	cat/cat_1003_32.jpg	0

```
In [17]: data['Label'].replace([0,1],['CATS','DOGS'],inplace=True)
data['Label'].value_counts().plot(figsize=(9,6),kind='bar', title = 'Dogs and Cats')
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3378cf3c10>
```



## Visualization of separation between the classes

Another way to represent images in a numerical format is to convert each image to a series of RGB pixels. For example, if we had a 32x32 px image, we'd convert it to a series of 1,024 RGB pixels. You could also explore other numerical measures like hue and saturation, but let's keep it simple for now.

```
In [18]: from PIL import Image
import numpy as np

#setup a standard image size; this will distort some images but will get every
#thing into the same shape
#STANDARD_SIZE = (32, 32)
def img_to_matrix(filename, verbose=False):
    """
    takes a filename and turns it into a numpy array of RGB pixels
    """
    img = Image.open(filename)

    #No need to resize the images, they are already 32x32
    #if verbose==True:
        # print "changing size from %s to %s" % (str(img.size), str(STANDARD_SI
ZE))
    #img = img.resize(STANDARD_SIZE)

    img = list(img.getdata())
    img = map(list, img)
    img = np.array(img)
    return img

def flatten_image(img):
    """
    takes in an (m, n) numpy array and flattens it
    into an array of shape (1, m * n)
    """
    s = img.shape[0] * img.shape[1]
    img_wide = img.reshape(1, s)
    return img_wide[0]
```

```
In [19]: # Read images from the directory and flatten the images.  
#To speed up the process we just use 100 images from each category (Dog and Cat)  
import os  
dog_img_dir = "/home/ubuntu/Data/dog_cat_32/train/dog/"  
dog_images = [dog_img_dir+f for f in os.listdir(dog_img_dir)]  
  
cat_img_dir = "/home/ubuntu/Data/dog_cat_32/train/cat/"  
cat_images = [cat_img_dir+f for f in os.listdir(cat_img_dir)]  
  
data = []  
labels = []  
  
for dog_image in dog_images[:300]:  
    img = img_to_matrix(dog_image)  
    img = flatten_image(img)  
    data.append(img)  
    labels.append("Dog")  
  
for cat_image in cat_images[:300]:  
    img = img_to_matrix(cat_image)  
    img = flatten_image(img)  
    data.append(img)  
    labels.append("Cats")  
  
data = np.array(data)
```

1,024 features is a lot to deal with for many algorithms, so we need to reduce the number of dimensions somehow. For this we can use an unsupervised learning technique called Randomized PCA to derive a smaller number of features from the raw pixel data.

Let's transform the dataset into just 2 components which we can easily plot in 2 dimensions.

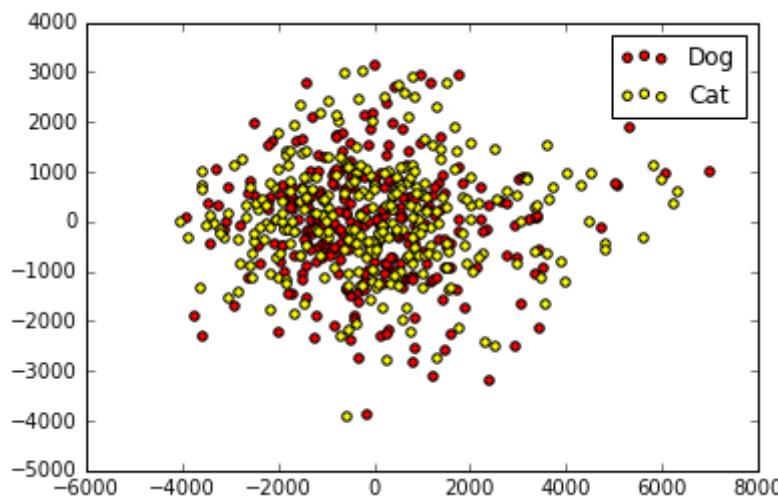
Reducing the data to 2 features maintains and allows us to visualize separation between the classes.

```
In [20]: from sklearn.decomposition import RandomizedPCA
import pylab as pl

pca = RandomizedPCA(n_components=2)
X = pca.fit_transform(data)
y = np.where(np.array(labels)=="Dog", 1, 0)

df = pd.DataFrame({"x": X[:, 0], "y": X[:, 1], "label":np.where(y==1, "Dog", "Cat")})

colors = ["red", "yellow"]
for label, color in zip(df['label'].unique(), colors):
    mask = df['label']==label
    pl.scatter(df[mask]['x'], df[mask]['y'], c=color, label=label)
pl.legend()
pl.show()
```



Based on the plot above, it looks like there's not a strong split b/w Dogs and Cats images. The images are all mixed up, therefore, we won't be able to use supervised classification techniques like K-Nearest Neighbors to classify Dogs and Cats images. For this reason, we will be using Deep Learning and Convolutional Neural Networks to classify the images.

# Algorithms and Techniques

We will use Python and a popular open source deep learning framework called [Caffe](http://caffe.berkeleyvision.org/) (<http://caffe.berkeleyvision.org/>) to build the image classifier.

[Caffe](http://caffe.berkeleyvision.org/) (<http://caffe.berkeleyvision.org/>) is a deep learning framework developed by the Berkeley Vision and Learning Center (BVLC). It is written in C++ and has Python and Matlab bindings.

- Caffe is fast due to its highly optimized C/CUDA backend which integrates GPU acceleration
- Caffe is still very accessible due to command line, Python and Matlab interfaces
- The wrapper interfaces make it very easy to integrate DNN training and deployment into larger data analytics workflows
- Caffe has a large open-source development community adding new features all the time
- Caffe has an associated model-zoo where researchers can upload trained models for others to fine tune or use for inference using their own data

Convolutional Neural Networks are a powerful artificial neural network technique.

These networks preserve the spatial structure of the problem and were developed for object recognition tasks such as handwritten digit recognition. They are popular because people are achieving state-of-the-art results on difficult computer vision and natural language processing tasks.

Given a dataset of gray scale images with the standardized size of  $32 \times 32$  pixels each, a traditional feedforward neural network would require 1024 input weights (plus one bias).

This is fair enough, but the flattening of the image matrix of pixels to a long vector of pixel values loses all of the spatial structure in the image. Unless all of the images are perfectly resized, the neural network will have great difficulty with the problem.

Convolutional Neural Networks expect and preserve the spatial relationship between pixels by learning internal feature representations using small squares of input data. Features are learned and used across the whole image, allowing for the objects in the images to be shifted or translated in the scene and still detectable by the network.

There are three types of layers in a Convolutional Neural Network:

1. Convolutional Layers.
2. Pooling Layers.
3. Fully-Connected Layers.

## ***Convolutional Layers***

Convolutional layers are comprised of filters and feature maps.

- The **filters** are the “neurons” of the layer. They have input weights and output a value.
- The **feature map** is the output of one filter applied to the previous layer.

## ***Pooling Layers***

The pooling layers down-sample the previous layers feature map.

- Pooling layers follow a sequence of one or more convolutional layers and are intended to consolidate the features learned and expressed in the previous layers feature map. As such, pooling may be considered a technique to compress or generalize feature representations and generally reduce the overfitting of the training data by the model.

### Fully Connected Layers

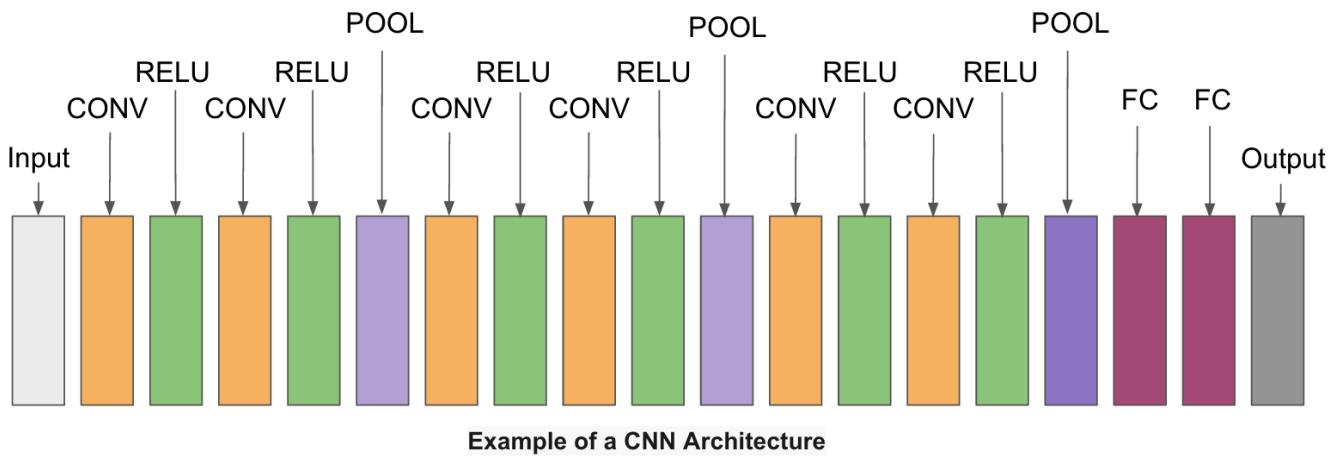
- Fully connected layers are the normal flat feed-forward neural network layer. These layers may have a non-linear activation function or a softmax activation in order to output probabilities of class predictions.

### Convolutional Neural Networks Architecture

The simplest architecture of a convolutional neural networks starts with an input layer (images) followed by a sequence of convolutional layers and pooling layers, and ends with fully-connected layers. The convolutional layers are usually followed by one layer of ReLU activation functions.

The convolutional, pooling and ReLU layers act as learnable features extractors, while the fully connected layers acts as a machine learning classifier. Furthermore, the early layers of the network encode generic patterns of the images, while later layers encode the details patterns of the images.

Note that only the convolutional layers and fully-connected layers have weights. These weights are learned in the training phase.



There are 4 steps in training a CNN using Caffe:

- Step 1 - Data preparation: In this step, we clean the images and store them in a format that can be used by Caffe. We use tools that comes with Caffe to move the training and validation images into **Imdb databases**. One for train (train.lmdb) and one for validation (val.lmdb)
- Step 2 - Model definition: In this step, we choose a CNN architecture and we define its parameters in a configuration file with extension **.prototxt**.
- Step 3 - Solver definition: The solver is responsible for model optimization. We define the solver parameters in a configuration file with extension **.prototxt**.
- Step 4 - Model training: We train the model by executing one Caffe command from the terminal. After training the model, we will get the trained model in a file with extension **.caffemodel**.

After the training phase, we will use the .caffemodel trained model to make predictions of new unseen data. We will write Python script to do this.

## Machine Setup

Images, videos, and other graphics are represented as matrices. GPUs, compared to CPUs, are faster at performing matrix operations and other advanced mathematical transformations. More detailed information regarding Hardware performance configurations can be found [here](http://caffe.berkeleyvision.org/performance_hardware.html) ([http://caffe.berkeleyvision.org/performance\\_hardware.html](http://caffe.berkeleyvision.org/performance_hardware.html))

For this project, [AWS \(https://aws.amazon.com/\)](https://aws.amazon.com/) EC2 instance of type g2.2xlarge is used. This instance has a high-performance NVIDIA GPU with 1,536 CUDA cores and 4GB of video memory, 15GB of RAM and 8 vCPUs. The machine costs \$0.65/hour.

You can skip the hassle of setting up deep learning frameworks from scratch by choosing an Amazon Machine Image (AMI) that comes pre-installed with the libraries and their dependencies. The Stanford class, [CS231n: Convolutional Neural Networks for Visual Recognition \(http://cs231n.stanford.edu/\)](http://cs231n.stanford.edu/), has provided a public AMI with these specs:

For step by step instruction on how to setup a AWS server with Jupyter, please follow this tutorial [Start deep learning with Jupyter notebooks in the cloud. \(http://efavdb.com/deep-learning-with-jupyter-on-aws/\)](http://efavdb.com/deep-learning-with-jupyter-on-aws/)

Let's execute the cell below to display information about the GPUs running on the server.

In [21]: !nvidia-smi

```

Wed Sep  7 18:03:25 2016
+-----+
| NVIDIA-SMI 352.63      Driver Version: 352.63      |
|-----+-----+-----+
| GPU  Name      Persistence-M| Bus-Id      Disp.A | Volatile Uncorr. ECC
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M.
|-----+-----+-----+-----+-----+-----+
|=|    0  GRID K520          Off   | 0000:00:03.0     Off  |                  N/A
|    | N/A   30C     P8    17W / 125W | 156MiB / 4095MiB | 0%       Default
|    |
+-----+-----+-----+
| Processes:                               GPU Memory
| GPU     PID  Type  Process name        Usage
|-----+-----+-----+-----+
|=|    0      1985  C    /home/ubuntu/anaconda2/bin/python  144MiB
|    |
+-----+-----+

```

## Benchmark

A similar problem was defined at [Kaggle dogs-vs-cats \(<https://www.kaggle.com/c/dogs-vs-cats>\)](https://www.kaggle.com/c/dogs-vs-cats) competition in 2014 where contestants were asked to write an algorithm to classify whether images contain either a dog or a cat. A Deep learning expert wins Kaggle Dogs vs Cats image competition with an almost perfect result of %99.

Although the data set used in this project is from ImageNet and not from Kaggle, I use Kaggle dogs vs cats competition accuracy scores as a benchmark to evaluate the performance of the models presented in the project.

## Methodology

## Data Preprocessing

Before we train a neural network using Caffe we will move the training and validation images into a database. The database allows Caffe to efficiently iterate over the image data during training. The training and validation datasets are independent subsets of the original image dataset. We will train the network using the training dataset and then test the networks performance using the validation dataset; that way we can be sure that the network performs well for images that it has never been trained on.

Images, videos, and other graphics are represented as matrices, and when you perform a certain operation, such as a camera rotation or a zoom in effect, all you are doing is applying some mathematical transformation to a matrix.

We complete each of these tasks using command line tools that come with Caffe. A number of useful utilities for data pre-processing, network training and network deployment can be found in the Caffe installation folder in `$CAFFE_ROOT/build/tools`

**(Note: you can ignore the "Failed to initialize libdc1394" warning messages in the output). libdc1394 is a library for controlling camera hardware.**

In [22]: %%bash

```
#Delete database files if they already exist.
rm -rf train_lmdb val_lmdb

#Setup environment variables
TOOLS=/home/ubuntu/caffe/build/tools
TRAIN_DATA_ROOT=/home/ubuntu/Data/dog_cat_32/train/
VAL_DATA_ROOT=/home/ubuntu/Data/dog_cat_32/val/

#Create the training database
$TOOLS/convert_imageset \
--shuffle \
$TRAIN_DATA_ROOT \
$TRAIN_DATA_ROOT/train.txt \
train_lmdb

#Create the validation database
$TOOLS/convert_imageset \
--shuffle \
$VAL_DATA_ROOT \
$VAL_DATA_ROOT/val.txt \
val_lmdb
```

```
I0907 18:03:28.444833 2063 convert_imageset.cpp:83] Shuffling data
I0907 18:03:28.580814 2063 convert_imageset.cpp:86] A total of 12904 images.
I0907 18:03:28.581188 2063 db_lmdb.cpp:38] Opened lmdb train_lmdb
I0907 18:03:28.667840 2063 convert_imageset.cpp:144] Processed 1000 files.
I0907 18:03:28.750430 2063 convert_imageset.cpp:144] Processed 2000 files.
I0907 18:03:28.839179 2063 convert_imageset.cpp:144] Processed 3000 files.
I0907 18:03:28.921123 2063 convert_imageset.cpp:144] Processed 4000 files.
I0907 18:03:29.011930 2063 convert_imageset.cpp:144] Processed 5000 files.
I0907 18:03:29.093874 2063 convert_imageset.cpp:144] Processed 6000 files.
I0907 18:03:29.176609 2063 convert_imageset.cpp:144] Processed 7000 files.
I0907 18:03:29.267001 2063 convert_imageset.cpp:144] Processed 8000 files.
I0907 18:03:29.353193 2063 convert_imageset.cpp:144] Processed 9000 files.
I0907 18:03:29.436241 2063 convert_imageset.cpp:144] Processed 10000 files.
I0907 18:03:29.519358 2063 convert_imageset.cpp:144] Processed 11000 files.
I0907 18:03:29.603425 2063 convert_imageset.cpp:144] Processed 12000 files.
I0907 18:03:29.685639 2063 convert_imageset.cpp:150] Processed 12904 files.
I0907 18:03:29.953532 2066 convert_imageset.cpp:83] Shuffling data
I0907 18:03:30.081560 2066 convert_imageset.cpp:86] A total of 1138 images.
I0907 18:03:30.081914 2066 db_lmdb.cpp:38] Opened lmdb val_lmdb
I0907 18:03:30.168824 2066 convert_imageset.cpp:144] Processed 1000 files.
I0907 18:03:30.186359 2066 convert_imageset.cpp:150] Processed 1138 files.
```

We also create a mean image from the training data. This is the image obtained by taking the mean value of each pixel across all of the training dataset images. We do this so that we can extract that mean image from each training and validation image before it is fed into the neural network. This is an important pre-processing and refinement step for achieving fast and effective training. It has the effect of removing the average brightness (intensity) of each point in the image so that the network learns about image content rather than illumination conditions.

Execute the cell below to create a mean image of the training data.

In [23]: %%bash

```
#Setup environment variables
TOOLS=/home/ubuntu/caffe/build/tools

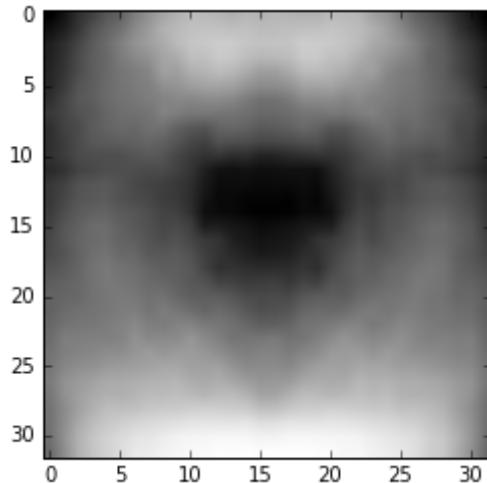
#Create the mean image database
$TOOLS/compute_image_mean train_lmdb mean.binaryproto
```

Run the cell below to see what the mean image looks like. Strangely, it looks a little like a mouse...

In [24]:

```
import caffe
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.cm as cm
%matplotlib inline

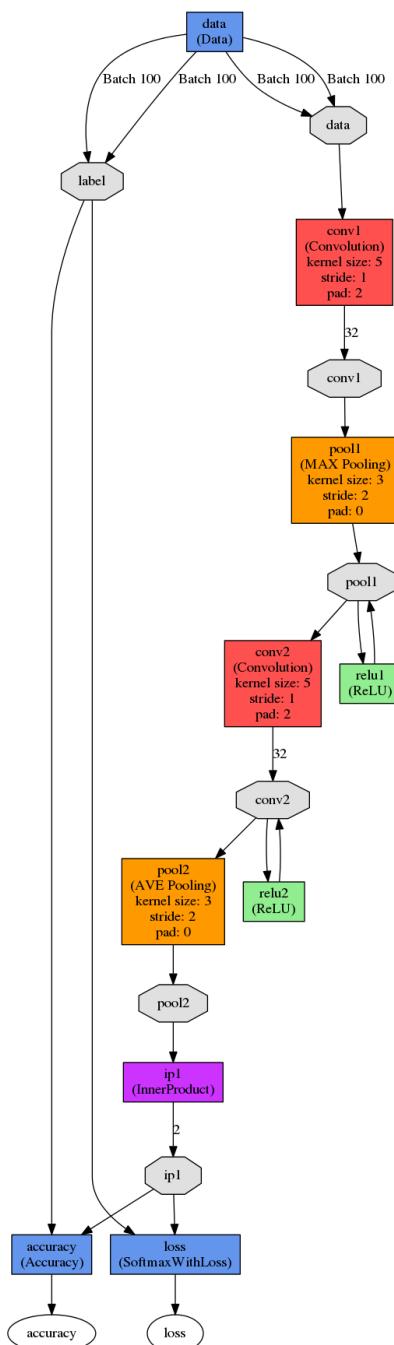
blob = caffe.proto.caffe_pb2.BlobProto()
data = open('/home/ubuntu/caffe/examples/Dog_Cat_Classification/mean.binaryproto','rb').read()
blob.ParseFromString(data)
arr = np.array(caffe.io.blobproto_to_array(blob))[0,:,:,:].mean(0)
plt.imshow(arr, cmap=cm.Greys_r)
plt.show()
```



## Implementation

We are now going to configure our network and start training. For this project we will borrow our network architecture from the cifar10 design that is provided with Caffe. You can find the original network in `$CAFFE_ROOT/examples/cifar10`. The only modification we will make initially is to change the final output softmax layer, which does the actual classification, to only have two classes for our dogs and cats problem rather than the ten classes needed for the CIFAR data.

Below we see an image of the cifar10 network architecture. As we can see the cifar10 network is a convolutional neural network (CNN) with three convolutional layers (each with pooling, ReLU activation and normalization) followed by a single fully-connected layer performing the final classification.



We see that the network input data has dimensions 100x3x32x32. This means that we have batches of 100 training images each with three-channels, i.e. it's color, and 32x32 pixels.

The first convolutional layer applies 32 5x5 filters. As the filters have size 5x5 we must pad the edges of the input images with zeros to ensure that the output of the layer has the same size.

After the convolutional layer max pooling is applied to reduce the size of the output images by half and a rectified linear unit (ReLU) activation function is applied.

Layers 2 is also a convolutional layer repeating this pattern but using average pooling.

The final layer is a fully connected layer with two neurons which together with the image labels feeds into a softmax layer which performs the classification and computes the classification loss. The two final neurons in the output layer correspond to our two classes, dogs and cats.

Caffe encodes deep neural network architectures like this in text files called prototxt files. You can get more information about the types of layers that can be defined in a network prototxt file [here](http://caffe.berkeleyvision.org/tutorial/layers.html) (<http://caffe.berkeleyvision.org/tutorial/layers.html>).

We now have our datasets and a network configuration we are nearly ready to train the network. The final thing that Caffe needs before we can train is a specification of the learning algorithm parameters. This specification is also made in a prototxt file but with a simpler structure. You can get more information about the range of parameters that can be set in a solver prototxt file [here](http://caffe.berkeleyvision.org/tutorial/solver.html) (<http://caffe.berkeleyvision.org/tutorial/solver.html>).

The network configuration and solver files are located at src directory.

We are now ready to train the network. Again, this is carried out using a command line tool that comes with Caffe, this time it is the binary `caffe` itself with the `train` option. Execute the cell below to begin training - it should take just under a minute to train - be sure to scroll down through the complete Caffe output.

```
In [8]: %%bash
#Set the location of the caffe tools folder
TOOLS=/home/ubuntu/caffe/build/tools
#Train the network
$TOOLS/caffe train -gpu 0 -solver src/solver.prototxt
```



```
I0907 17:49:10.599663 2018 upgrade_proto.cpp:990] Attempting to upgrade input file specified using deprecated 'solver_type' field (enum)': src/solver.prototxt
I0907 17:49:10.599895 2018 upgrade_proto.cpp:997] Successfully upgraded file specified using deprecated 'solver_type' field (enum) to 'type' field (string).
W0907 17:49:10.599907 2018 upgrade_proto.cpp:999] Note that future Caffe releases will only support 'type' field (string) for a solver's type.
I0907 17:49:10.600075 2018 caffe.cpp:185] Using GPUs 0
I0907 17:49:10.734774 2018 solver.cpp:48] Initializing solver from parameters:
test_iter: 100
test_interval: 250
base_lr: 0.01
display: 20
max_iter: 750
lr_policy: "step"
gamma: 0.1
momentum: 0.9
weight_decay: 0.0005
stepsize: 500
snapshot: 250
snapshot_prefix: "checkpoints/snapshot"
solver_mode: GPU
device_id: 0
random_seed: 1234
net: "src/train_val.prototxt"
type: "SGD"
I0907 17:49:10.734944 2018 solver.cpp:91] Creating training net from net file: src/train_val.prototxt
I0907 17:49:10.735399 2018 net.cpp:322] The NetState phase (0) differed from the phase (1) specified by a rule in layer data
I0907 17:49:10.735427 2018 net.cpp:322] The NetState phase (0) differed from the phase (1) specified by a rule in layer accuracy
I0907 17:49:10.735538 2018 net.cpp:49] Initializing net from parameters:
state {
    phase: TRAIN
}
layer {
    name: "data"
    type: "Data"
    top: "data"
    top: "label"
    include {
        phase: TRAIN
    }
    transform_param {
        mirror: true
        crop_size: 32
        mean_file: "mean.binaryproto"
    }
    data_param {
        source: "train_lmdb"
        batch_size: 100
        backend: LMDB
    }
}
```

```
layer {
    name: "conv1"
    type: "Convolution"
    bottom: "data"
    top: "conv1"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 32
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.0001
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "pool1"
    type: "Pooling"
    bottom: "conv1"
    top: "pool1"
    pooling_param {
        pool: MAX
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "relu1"
    type: "ReLU"
    bottom: "pool1"
    top: "pool1"
}
layer {
    name: "conv2"
    type: "Convolution"
    bottom: "pool1"
    top: "conv2"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 32
        pad: 2
        kernel_size: 5
    }
}
```

```
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "relu2"
    type: "ReLU"
    bottom: "conv2"
    top: "conv2"
}
layer {
    name: "pool2"
    type: "Pooling"
    bottom: "conv2"
    top: "pool2"
    pooling_param {
        pool: AVE
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "ip1"
    type: "InnerProduct"
    bottom: "pool2"
    top: "ip1"
    param {
        lr_mult: 1
        decay_mult: 250
    }
    param {
        lr_mult: 2
        decay_mult: 0
    }
    inner_product_param {
        num_output: 2
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "loss"
    type: "SoftmaxWithLoss"
    bottom: "ip1"
    bottom: "label"
    top: "loss"
```

```
}
```

I0907 17:49:10.735671 2018 layer\_factory.hpp:77] Creating layer data  
I0907 17:49:10.735867 2018 net.cpp:106] Creating Layer data  
I0907 17:49:10.735889 2018 net.cpp:411] data -> data  
I0907 17:49:10.735931 2018 net.cpp:411] data -> label  
I0907 17:49:10.735960 2018 data\_transformer.cpp:25] Loading mean file from:  
mean.binaryproto  
I0907 17:49:10.736557 2021 db\_lmdb.cpp:38] Opened lmdb train\_lmdb  
I0907 17:49:10.747678 2018 data\_layer.cpp:41] output data size: 100,3,32,32  
I0907 17:49:10.750541 2018 net.cpp:150] Setting up data  
I0907 17:49:10.750619 2018 net.cpp:157] Top shape: 100 3 32 32 (307200)  
I0907 17:49:10.750636 2018 net.cpp:157] Top shape: 100 (100)  
I0907 17:49:10.750641 2018 net.cpp:165] Memory required for data: 1229200  
I0907 17:49:10.750655 2018 layer\_factory.hpp:77] Creating layer conv1  
I0907 17:49:10.750687 2018 net.cpp:106] Creating Layer conv1  
I0907 17:49:10.750704 2018 net.cpp:454] conv1 <- data  
I0907 17:49:10.750721 2018 net.cpp:411] conv1 -> conv1  
I0907 17:49:10.869427 2018 net.cpp:150] Setting up conv1  
I0907 17:49:10.869477 2018 net.cpp:157] Top shape: 100 32 32 32 (3276800)  
I0907 17:49:10.869514 2018 net.cpp:165] Memory required for data: 14336400  
I0907 17:49:10.869545 2018 layer\_factory.hpp:77] Creating layer pool1  
I0907 17:49:10.869575 2018 net.cpp:106] Creating Layer pool1  
I0907 17:49:10.869588 2018 net.cpp:454] pool1 <- conv1  
I0907 17:49:10.869598 2018 net.cpp:411] pool1 -> pool1  
I0907 17:49:10.869664 2018 net.cpp:150] Setting up pool1  
I0907 17:49:10.869681 2018 net.cpp:157] Top shape: 100 32 16 16 (819200)  
I0907 17:49:10.869686 2018 net.cpp:165] Memory required for data: 17613200  
I0907 17:49:10.869701 2018 layer\_factory.hpp:77] Creating layer relu1  
I0907 17:49:10.869710 2018 net.cpp:106] Creating Layer relu1  
I0907 17:49:10.869720 2018 net.cpp:454] relu1 <- pool1  
I0907 17:49:10.869729 2018 net.cpp:397] relu1 -> pool1 (in-place)  
I0907 17:49:10.869884 2018 net.cpp:150] Setting up relu1  
I0907 17:49:10.869904 2018 net.cpp:157] Top shape: 100 32 16 16 (819200)  
I0907 17:49:10.869909 2018 net.cpp:165] Memory required for data: 20890000  
I0907 17:49:10.869915 2018 layer\_factory.hpp:77] Creating layer conv2  
I0907 17:49:10.869933 2018 net.cpp:106] Creating Layer conv2  
I0907 17:49:10.869943 2018 net.cpp:454] conv2 <- pool1  
I0907 17:49:10.869952 2018 net.cpp:411] conv2 -> conv2  
I0907 17:49:10.872194 2018 net.cpp:150] Setting up conv2  
I0907 17:49:10.872220 2018 net.cpp:157] Top shape: 100 32 16 16 (819200)  
I0907 17:49:10.872225 2018 net.cpp:165] Memory required for data: 24166800  
I0907 17:49:10.872241 2018 layer\_factory.hpp:77] Creating layer relu2  
I0907 17:49:10.872254 2018 net.cpp:106] Creating Layer relu2  
I0907 17:49:10.872261 2018 net.cpp:454] relu2 <- conv2  
I0907 17:49:10.872272 2018 net.cpp:397] relu2 -> conv2 (in-place)  
I0907 17:49:10.872520 2018 net.cpp:150] Setting up relu2  
I0907 17:49:10.872550 2018 net.cpp:157] Top shape: 100 32 16 16 (819200)  
I0907 17:49:10.872555 2018 net.cpp:165] Memory required for data: 27443600  
I0907 17:49:10.872566 2018 layer\_factory.hpp:77] Creating layer pool2  
I0907 17:49:10.872580 2018 net.cpp:106] Creating Layer pool2  
I0907 17:49:10.872591 2018 net.cpp:454] pool2 <- conv2  
I0907 17:49:10.872609 2018 net.cpp:411] pool2 -> pool12  
I0907 17:49:10.872788 2018 net.cpp:150] Setting up pool2  
I0907 17:49:10.872807 2018 net.cpp:157] Top shape: 100 32 8 8 (204800)  
I0907 17:49:10.872812 2018 net.cpp:165] Memory required for data: 28262800  
I0907 17:49:10.872819 2018 layer\_factory.hpp:77] Creating layer ip1  
I0907 17:49:10.872833 2018 net.cpp:106] Creating Layer ip1

```
I0907 17:49:10.872844 2018 net.cpp:454] ip1 <- pool2
I0907 17:49:10.872856 2018 net.cpp:411] ip1 -> ip1
I0907 17:49:10.873124 2018 net.cpp:150] Setting up ip1
I0907 17:49:10.873142 2018 net.cpp:157] Top shape: 100 2 (200)
I0907 17:49:10.873147 2018 net.cpp:165] Memory required for data: 28263600
I0907 17:49:10.873160 2018 layer_factory.hpp:77] Creating layer loss
I0907 17:49:10.873183 2018 net.cpp:106] Creating Layer loss
I0907 17:49:10.873194 2018 net.cpp:454] loss <- ip1
I0907 17:49:10.873201 2018 net.cpp:454] loss <- label
I0907 17:49:10.873214 2018 net.cpp:411] loss -> loss
I0907 17:49:10.873237 2018 layer_factory.hpp:77] Creating layer loss
I0907 17:49:10.873580 2018 net.cpp:150] Setting up loss
I0907 17:49:10.873600 2018 net.cpp:157] Top shape: (1)
I0907 17:49:10.873606 2018 net.cpp:160] with loss weight 1
I0907 17:49:10.873637 2018 net.cpp:165] Memory required for data: 28263604
I0907 17:49:10.873643 2018 net.cpp:226] loss needs backward computation.
I0907 17:49:10.873649 2018 net.cpp:226] ip1 needs backward computation.
I0907 17:49:10.873656 2018 net.cpp:226] pool2 needs backward computation.
I0907 17:49:10.873659 2018 net.cpp:226] relu2 needs backward computation.
I0907 17:49:10.873666 2018 net.cpp:226] conv2 needs backward computation.
I0907 17:49:10.873669 2018 net.cpp:226] relu1 needs backward computation.
I0907 17:49:10.873675 2018 net.cpp:226] pool1 needs backward computation.
I0907 17:49:10.873679 2018 net.cpp:226] conv1 needs backward computation.
I0907 17:49:10.873685 2018 net.cpp:228] data does not need backward computation.

I0907 17:49:10.873709 2018 net.cpp:270] This network produces output loss
I0907 17:49:10.873724 2018 net.cpp:283] Network initialization done.
I0907 17:49:10.874142 2018 solver.cpp:181] Creating test net (#0) specified by net file: src/train_val.prototxt
I0907 17:49:10.874191 2018 net.cpp:322] The NetState phase (1) differed from the phase (0) specified by a rule in layer data
I0907 17:49:10.874311 2018 net.cpp:49] Initializing net from parameters:
state {
    phase: TEST
}
layer {
    name: "data"
    type: "Data"
    top: "data"
    top: "label"
    include {
        phase: TEST
    }
    transform_param {
        crop_size: 32
        mean_file: "mean.binaryproto"
    }
    data_param {
        source: "val_lmdb"
        batch_size: 100
        backend: LMDB
    }
}
layer {
    name: "conv1"
    type: "Convolution"
    bottom: "data"
```

```
top: "conv1"
param {
    lr_mult: 1
}
param {
    lr_mult: 2
}
convolution_param {
    num_output: 32
    pad: 2
    kernel_size: 5
    stride: 1
    weight_filler {
        type: "gaussian"
        std: 0.0001
    }
    bias_filler {
        type: "constant"
    }
}
layer {
    name: "pool1"
    type: "Pooling"
    bottom: "conv1"
    top: "pool1"
    pooling_param {
        pool: MAX
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "relu1"
    type: "ReLU"
    bottom: "pool1"
    top: "pool1"
}
layer {
    name: "conv2"
    type: "Convolution"
    bottom: "pool1"
    top: "conv2"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 32
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
    }
}
```

```
        }
        bias_filler {
            type: "constant"
        }
    }
layer {
    name: "relu2"
    type: "ReLU"
    bottom: "conv2"
    top: "conv2"
}
layer {
    name: "pool2"
    type: "Pooling"
    bottom: "conv2"
    top: "pool2"
    pooling_param {
        pool: AVE
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "ip1"
    type: "InnerProduct"
    bottom: "pool2"
    top: "ip1"
    param {
        lr_mult: 1
        decay_mult: 250
    }
    param {
        lr_mult: 2
        decay_mult: 0
    }
    inner_product_param {
        num_output: 2
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "loss"
    type: "SoftmaxWithLoss"
    bottom: "ip1"
    bottom: "label"
    top: "loss"
}
layer {
    name: "accuracy"
    type: "Accuracy"
```

```
bottom: "ip1"
bottom: "label"
top: "accuracy"
include {
    phase: TEST
}
}
I0907 17:49:10.874408 2018 layer_factory.hpp:77] Creating layer data
I0907 17:49:10.874550 2018 net.cpp:106] Creating Layer data
I0907 17:49:10.874567 2018 net.cpp:411] data -> data
I0907 17:49:10.874579 2018 net.cpp:411] data -> label
I0907 17:49:10.874595 2018 data_transformer.cpp:25] Loading mean file from:
mean.binaryproto
I0907 17:49:10.875272 2023 db_lmdb.cpp:38] Opened lmdb val_lmdb
I0907 17:49:10.875392 2018 data_layer.cpp:41] output data size: 100,3,32,32
I0907 17:49:10.878393 2018 net.cpp:150] Setting up data
I0907 17:49:10.878414 2018 net.cpp:157] Top shape: 100 3 32 32 (307200)
I0907 17:49:10.878422 2018 net.cpp:157] Top shape: 100 (100)
I0907 17:49:10.878433 2018 net.cpp:165] Memory required for data: 1229200
I0907 17:49:10.878439 2018 layer_factory.hpp:77] Creating layer label_data_1
_split
I0907 17:49:10.878448 2018 net.cpp:106] Creating Layer label_data_1_split
I0907 17:49:10.878459 2018 net.cpp:454] label_data_1_split <- label
I0907 17:49:10.878466 2018 net.cpp:411] label_data_1_split -> label_data_1_s
plit_0
I0907 17:49:10.878481 2018 net.cpp:411] label_data_1_split -> label_data_1_s
plit_1
I0907 17:49:10.878564 2018 net.cpp:150] Setting up label_data_1_split
I0907 17:49:10.878581 2018 net.cpp:157] Top shape: 100 (100)
I0907 17:49:10.878587 2018 net.cpp:157] Top shape: 100 (100)
I0907 17:49:10.878597 2018 net.cpp:165] Memory required for data: 1230000
I0907 17:49:10.878602 2018 layer_factory.hpp:77] Creating layer conv1
I0907 17:49:10.878646 2018 net.cpp:106] Creating Layer conv1
I0907 17:49:10.878660 2018 net.cpp:454] conv1 <- data
I0907 17:49:10.878669 2018 net.cpp:411] conv1 -> conv1
I0907 17:49:10.879997 2018 net.cpp:150] Setting up conv1
I0907 17:49:10.880022 2018 net.cpp:157] Top shape: 100 32 32 32 (3276800)
I0907 17:49:10.880028 2018 net.cpp:165] Memory required for data: 14337200
I0907 17:49:10.880043 2018 layer_factory.hpp:77] Creating layer pool1
I0907 17:49:10.880056 2018 net.cpp:106] Creating Layer pool1
I0907 17:49:10.880069 2018 net.cpp:454] pool1 <- conv1
I0907 17:49:10.880076 2018 net.cpp:411] pool1 -> pool1
I0907 17:49:10.880125 2018 net.cpp:150] Setting up pool1
I0907 17:49:10.880141 2018 net.cpp:157] Top shape: 100 32 16 16 (819200)
I0907 17:49:10.880146 2018 net.cpp:165] Memory required for data: 17614000
I0907 17:49:10.880153 2018 layer_factory.hpp:77] Creating layer relu1
I0907 17:49:10.880162 2018 net.cpp:106] Creating Layer relu1
I0907 17:49:10.880172 2018 net.cpp:454] relu1 <- pool1
I0907 17:49:10.880178 2018 net.cpp:397] relu1 -> pool1 (in-place)
I0907 17:49:10.880343 2018 net.cpp:150] Setting up relu1
I0907 17:49:10.880365 2018 net.cpp:157] Top shape: 100 32 16 16 (819200)
I0907 17:49:10.880370 2018 net.cpp:165] Memory required for data: 20890800
I0907 17:49:10.880378 2018 layer_factory.hpp:77] Creating layer conv2
I0907 17:49:10.880393 2018 net.cpp:106] Creating Layer conv2
I0907 17:49:10.880404 2018 net.cpp:454] conv2 <- pool1
I0907 17:49:10.880422 2018 net.cpp:411] conv2 -> conv2
I0907 17:49:10.882477 2018 net.cpp:150] Setting up conv2
```

```
I0907 17:49:10.882499 2018 net.cpp:157] Top shape: 100 32 16 16 (819200)
I0907 17:49:10.882505 2018 net.cpp:165] Memory required for data: 24167600
I0907 17:49:10.882519 2018 layer_factory.hpp:77] Creating layer relu2
I0907 17:49:10.882536 2018 net.cpp:106] Creating Layer relu2
I0907 17:49:10.882545 2018 net.cpp:454] relu2 <- conv2
I0907 17:49:10.882552 2018 net.cpp:397] relu2 -> conv2 (in-place)
I0907 17:49:10.882721 2018 net.cpp:150] Setting up relu2
I0907 17:49:10.882740 2018 net.cpp:157] Top shape: 100 32 16 16 (819200)
I0907 17:49:10.882745 2018 net.cpp:165] Memory required for data: 27444400
I0907 17:49:10.882752 2018 layer_factory.hpp:77] Creating layer pool2
I0907 17:49:10.882761 2018 net.cpp:106] Creating Layer pool2
I0907 17:49:10.882771 2018 net.cpp:454] pool2 <- conv2
I0907 17:49:10.882781 2018 net.cpp:411] pool2 -> pool2
I0907 17:49:10.883066 2018 net.cpp:150] Setting up pool2
I0907 17:49:10.883086 2018 net.cpp:157] Top shape: 100 32 8 8 (204800)
I0907 17:49:10.883091 2018 net.cpp:165] Memory required for data: 28263600
I0907 17:49:10.883098 2018 layer_factory.hpp:77] Creating layer ip1
I0907 17:49:10.883110 2018 net.cpp:106] Creating Layer ip1
I0907 17:49:10.883119 2018 net.cpp:454] ip1 <- pool2
I0907 17:49:10.883134 2018 net.cpp:411] ip1 -> ip1
I0907 17:49:10.883389 2018 net.cpp:150] Setting up ip1
I0907 17:49:10.883407 2018 net.cpp:157] Top shape: 100 2 (200)
I0907 17:49:10.883412 2018 net.cpp:165] Memory required for data: 28264400
I0907 17:49:10.883425 2018 layer_factory.hpp:77] Creating layer ip1_ip1_0_sp
lit
I0907 17:49:10.883440 2018 net.cpp:106] Creating Layer ip1_ip1_0_split
I0907 17:49:10.883445 2018 net.cpp:454] ip1_ip1_0_split <- ip1
I0907 17:49:10.883457 2018 net.cpp:411] ip1_ip1_0_split -> ip1_ip1_0_split_0
I0907 17:49:10.883471 2018 net.cpp:411] ip1_ip1_0_split -> ip1_ip1_0_split_1
I0907 17:49:10.883517 2018 net.cpp:150] Setting up ip1_ip1_0_split
I0907 17:49:10.883533 2018 net.cpp:157] Top shape: 100 2 (200)
I0907 17:49:10.883539 2018 net.cpp:157] Top shape: 100 2 (200)
I0907 17:49:10.883544 2018 net.cpp:165] Memory required for data: 28266000
I0907 17:49:10.883550 2018 layer_factory.hpp:77] Creating layer loss
I0907 17:49:10.883558 2018 net.cpp:106] Creating Layer loss
I0907 17:49:10.883569 2018 net.cpp:454] loss <- ip1_ip1_0_split_0
I0907 17:49:10.883575 2018 net.cpp:454] loss <- label_data_1_split_0
I0907 17:49:10.883592 2018 net.cpp:411] loss -> loss
I0907 17:49:10.883626 2018 layer_factory.hpp:77] Creating layer loss
I0907 17:49:10.883978 2018 net.cpp:150] Setting up loss
I0907 17:49:10.884004 2018 net.cpp:157] Top shape: (1)
I0907 17:49:10.884016 2018 net.cpp:160] with loss weight 1
I0907 17:49:10.884065 2018 net.cpp:165] Memory required for data: 28266004
I0907 17:49:10.884078 2018 layer_factory.hpp:77] Creating layer accuracy
I0907 17:49:10.884096 2018 net.cpp:106] Creating Layer accuracy
I0907 17:49:10.884116 2018 net.cpp:454] accuracy <- ip1_ip1_0_split_1
I0907 17:49:10.884130 2018 net.cpp:454] accuracy <- label_data_1_split_1
I0907 17:49:10.884152 2018 net.cpp:411] accuracy -> accuracy
I0907 17:49:10.884187 2018 net.cpp:150] Setting up accuracy
I0907 17:49:10.884209 2018 net.cpp:157] Top shape: (1)
I0907 17:49:10.884219 2018 net.cpp:165] Memory required for data: 28266008
I0907 17:49:10.884232 2018 net.cpp:228] accuracy does not need backward computation.
I0907 17:49:10.884244 2018 net.cpp:226] loss needs backward computation.
I0907 17:49:10.884263 2018 net.cpp:226] ip1_ip1_0_split needs backward computation.
I0907 17:49:10.884274 2018 net.cpp:226] ip1 needs backward computation.
```

```
I0907 17:49:10.884290 2018 net.cpp:226] pool2 needs backward computation.
I0907 17:49:10.884300 2018 net.cpp:226] relu2 needs backward computation.
I0907 17:49:10.884312 2018 net.cpp:226] conv2 needs backward computation.
I0907 17:49:10.884322 2018 net.cpp:226] relu1 needs backward computation.
I0907 17:49:10.884340 2018 net.cpp:226] pool1 needs backward computation.
I0907 17:49:10.884349 2018 net.cpp:226] conv1 needs backward computation.
I0907 17:49:10.884364 2018 net.cpp:228] label_data_1_split does not need backward computation.
I0907 17:49:10.884374 2018 net.cpp:228] data does not need backward computation.
I0907 17:49:10.884390 2018 net.cpp:270] This network produces output accuracy
I0907 17:49:10.884402 2018 net.cpp:270] This network produces output loss
I0907 17:49:10.884426 2018 net.cpp:283] Network initialization done.
I0907 17:49:10.884502 2018 solver.cpp:60] Solver scaffolding done.
I0907 17:49:10.884750 2018 caffe.cpp:213] Starting Optimization
I0907 17:49:10.884773 2018 solver.cpp:280] Solving
I0907 17:49:10.884781 2018 solver.cpp:281] Learning Rate Policy: step
I0907 17:49:10.885170 2018 solver.cpp:338] Iteration 0, Testing net (#0)
I0907 17:49:11.375325 2018 solver.cpp:406] Test net output #0: accuracy
= 0.5206
I0907 17:49:11.375389 2018 solver.cpp:406] Test net output #1: loss = 0.
693017 (* 1 = 0.693017 loss)
I0907 17:49:11.381693 2018 solver.cpp:229] Iteration 0, loss = 0.693084
I0907 17:49:11.381729 2018 solver.cpp:245] Train net output #0: loss =
0.693084 (* 1 = 0.693084 loss)
I0907 17:49:11.381758 2018 sgd_solver.cpp:106] Iteration 0, lr = 0.01
I0907 17:49:11.658129 2018 solver.cpp:229] Iteration 20, loss = 0.696418
I0907 17:49:11.658182 2018 solver.cpp:245] Train net output #0: loss =
0.696418 (* 1 = 0.696418 loss)
I0907 17:49:11.658197 2018 sgd_solver.cpp:106] Iteration 20, lr = 0.01
I0907 17:49:11.934586 2018 solver.cpp:229] Iteration 40, loss = 0.702814
I0907 17:49:11.934641 2018 solver.cpp:245] Train net output #0: loss =
0.702814 (* 1 = 0.702814 loss)
I0907 17:49:11.934656 2018 sgd_solver.cpp:106] Iteration 40, lr = 0.01
I0907 17:49:12.211274 2018 solver.cpp:229] Iteration 60, loss = 0.693217
I0907 17:49:12.211331 2018 solver.cpp:245] Train net output #0: loss =
0.693217 (* 1 = 0.693217 loss)
I0907 17:49:12.211347 2018 sgd_solver.cpp:106] Iteration 60, lr = 0.01
I0907 17:49:12.487753 2018 solver.cpp:229] Iteration 80, loss = 0.692173
I0907 17:49:12.487807 2018 solver.cpp:245] Train net output #0: loss =
0.692173 (* 1 = 0.692173 loss)
I0907 17:49:12.487823 2018 sgd_solver.cpp:106] Iteration 80, lr = 0.01
I0907 17:49:12.764281 2018 solver.cpp:229] Iteration 100, loss = 0.693186
I0907 17:49:12.764334 2018 solver.cpp:245] Train net output #0: loss =
0.693186 (* 1 = 0.693186 loss)
I0907 17:49:12.764351 2018 sgd_solver.cpp:106] Iteration 100, lr = 0.01
I0907 17:49:13.040750 2018 solver.cpp:229] Iteration 120, loss = 0.693745
I0907 17:49:13.040812 2018 solver.cpp:245] Train net output #0: loss =
0.693745 (* 1 = 0.693745 loss)
I0907 17:49:13.040828 2018 sgd_solver.cpp:106] Iteration 120, lr = 0.01
I0907 17:49:13.318050 2018 solver.cpp:229] Iteration 140, loss = 0.688764
I0907 17:49:13.318100 2018 solver.cpp:245] Train net output #0: loss =
0.688764 (* 1 = 0.688764 loss)
I0907 17:49:13.318116 2018 sgd_solver.cpp:106] Iteration 140, lr = 0.01
I0907 17:49:13.595294 2018 solver.cpp:229] Iteration 160, loss = 0.694766
I0907 17:49:13.595350 2018 solver.cpp:245] Train net output #0: loss =
```

```

0.694766 (* 1 = 0.694766 loss)
I0907 17:49:13.595366 2018 sgd_solver.cpp:106] Iteration 160, lr = 0.01
I0907 17:49:13.872601 2018 solver.cpp:229] Iteration 180, loss = 0.68917
I0907 17:49:13.872658 2018 solver.cpp:245]      Train net output #0: loss =
0.68917 (* 1 = 0.68917 loss)
I0907 17:49:13.872673 2018 sgd_solver.cpp:106] Iteration 180, lr = 0.01
I0907 17:49:14.149952 2018 solver.cpp:229] Iteration 200, loss = 0.694095
I0907 17:49:14.150007 2018 solver.cpp:245]      Train net output #0: loss =
0.694095 (* 1 = 0.694095 loss)
I0907 17:49:14.150022 2018 sgd_solver.cpp:106] Iteration 200, lr = 0.01
I0907 17:49:14.427142 2018 solver.cpp:229] Iteration 220, loss = 0.694301
I0907 17:49:14.427196 2018 solver.cpp:245]      Train net output #0: loss =
0.694301 (* 1 = 0.694301 loss)
I0907 17:49:14.427211 2018 sgd_solver.cpp:106] Iteration 220, lr = 0.01
I0907 17:49:14.704120 2018 solver.cpp:229] Iteration 240, loss = 0.691377
I0907 17:49:14.704172 2018 solver.cpp:245]      Train net output #0: loss =
0.691377 (* 1 = 0.691377 loss)
I0907 17:49:14.704187 2018 sgd_solver.cpp:106] Iteration 240, lr = 0.01
I0907 17:49:14.828528 2018 solver.cpp:456] Snapshotting to binary proto file
checkpoints/snapshot_iter_250.caffemodel
I0907 17:49:14.838659 2018 sgd_solver.cpp:273] Snapshotting solver state to
binary proto file checkpoints/snapshot_iter_250.solverstate
I0907 17:49:14.839365 2018 solver.cpp:338] Iteration 250, Testing net (#0)
I0907 17:49:15.329273 2018 solver.cpp:406]      Test net output #0: accuracy
= 0.5062
I0907 17:49:15.329335 2018 solver.cpp:406]      Test net output #1: loss = 0.
694212 (* 1 = 0.694212 loss)
I0907 17:49:15.472729 2018 solver.cpp:229] Iteration 260, loss = 0.696516
I0907 17:49:15.472792 2018 solver.cpp:245]      Train net output #0: loss =
0.696516 (* 1 = 0.696516 loss)
I0907 17:49:15.472807 2018 sgd_solver.cpp:106] Iteration 260, lr = 0.01
I0907 17:49:15.749964 2018 solver.cpp:229] Iteration 280, loss = 0.698534
I0907 17:49:15.750020 2018 solver.cpp:245]      Train net output #0: loss =
0.698534 (* 1 = 0.698534 loss)
I0907 17:49:15.750036 2018 sgd_solver.cpp:106] Iteration 280, lr = 0.01
I0907 17:49:16.026904 2018 solver.cpp:229] Iteration 300, loss = 0.690629
I0907 17:49:16.026959 2018 solver.cpp:245]      Train net output #0: loss =
0.690629 (* 1 = 0.690629 loss)
I0907 17:49:16.026974 2018 sgd_solver.cpp:106] Iteration 300, lr = 0.01
I0907 17:49:16.304231 2018 solver.cpp:229] Iteration 320, loss = 0.691888
I0907 17:49:16.304287 2018 solver.cpp:245]      Train net output #0: loss =
0.691888 (* 1 = 0.691888 loss)
I0907 17:49:16.304303 2018 sgd_solver.cpp:106] Iteration 320, lr = 0.01
I0907 17:49:16.580847 2018 solver.cpp:229] Iteration 340, loss = 0.69386
I0907 17:49:16.580900 2018 solver.cpp:245]      Train net output #0: loss =
0.69386 (* 1 = 0.69386 loss)
I0907 17:49:16.580915 2018 sgd_solver.cpp:106] Iteration 340, lr = 0.01
I0907 17:49:16.857722 2018 solver.cpp:229] Iteration 360, loss = 0.690541
I0907 17:49:16.857780 2018 solver.cpp:245]      Train net output #0: loss =
0.69054 (* 1 = 0.69054 loss)
I0907 17:49:16.857796 2018 sgd_solver.cpp:106] Iteration 360, lr = 0.01
I0907 17:49:17.134531 2018 solver.cpp:229] Iteration 380, loss = 0.699985
I0907 17:49:17.134583 2018 solver.cpp:245]      Train net output #0: loss =
0.699985 (* 1 = 0.699985 loss)
I0907 17:49:17.134598 2018 sgd_solver.cpp:106] Iteration 380, lr = 0.01
I0907 17:49:17.411090 2018 solver.cpp:229] Iteration 400, loss = 0.698098
I0907 17:49:17.411149 2018 solver.cpp:245]      Train net output #0: loss =

```

```

0.698098 (* 1 = 0.698098 loss)
I0907 17:49:17.411164 2018 sgd_solver.cpp:106] Iteration 400, lr = 0.01
I0907 17:49:17.687440 2018 solver.cpp:229] Iteration 420, loss = 0.695981
I0907 17:49:17.687492 2018 solver.cpp:245] Train net output #0: loss =
0.695981 (* 1 = 0.695981 loss)
I0907 17:49:17.687508 2018 sgd_solver.cpp:106] Iteration 420, lr = 0.01
I0907 17:49:17.963994 2018 solver.cpp:229] Iteration 440, loss = 0.693345
I0907 17:49:17.964051 2018 solver.cpp:245] Train net output #0: loss =
0.693345 (* 1 = 0.693345 loss)
I0907 17:49:17.964074 2018 sgd_solver.cpp:106] Iteration 440, lr = 0.01
I0907 17:49:18.241008 2018 solver.cpp:229] Iteration 460, loss = 0.6923
I0907 17:49:18.241071 2018 solver.cpp:245] Train net output #0: loss =
0.6923 (* 1 = 0.6923 loss)
I0907 17:49:18.241087 2018 sgd_solver.cpp:106] Iteration 460, lr = 0.01
I0907 17:49:18.517881 2018 solver.cpp:229] Iteration 480, loss = 0.691581
I0907 17:49:18.517935 2018 solver.cpp:245] Train net output #0: loss =
0.691581 (* 1 = 0.691581 loss)
I0907 17:49:18.517951 2018 sgd_solver.cpp:106] Iteration 480, lr = 0.01
I0907 17:49:18.780949 2018 solver.cpp:456] Snapshotting to binary proto file
checkpoints/snapshot_iter_500.caffemodel
I0907 17:49:18.790834 2018 sgd_solver.cpp:273] Snapshotting solver state to
binary proto file checkpoints/snapshot_iter_500.solverstate
I0907 17:49:18.791682 2018 solver.cpp:338] Iteration 500, Testing net (#0)
I0907 17:49:19.280625 2018 solver.cpp:406] Test net output #0: accuracy
= 0.5071
I0907 17:49:19.280684 2018 solver.cpp:406] Test net output #1: loss = 0.
693794 (* 1 = 0.693794 loss)
I0907 17:49:19.285815 2018 solver.cpp:229] Iteration 500, loss = 0.694545
I0907 17:49:19.285861 2018 solver.cpp:245] Train net output #0: loss =
0.694545 (* 1 = 0.694545 loss)
I0907 17:49:19.285878 2018 sgd_solver.cpp:106] Iteration 500, lr = 0.001
I0907 17:49:19.563009 2018 solver.cpp:229] Iteration 520, loss = 0.688978
I0907 17:49:19.563076 2018 solver.cpp:245] Train net output #0: loss =
0.688978 (* 1 = 0.688978 loss)
I0907 17:49:19.563091 2018 sgd_solver.cpp:106] Iteration 520, lr = 0.001
I0907 17:49:19.840301 2018 solver.cpp:229] Iteration 540, loss = 0.689023
I0907 17:49:19.840355 2018 solver.cpp:245] Train net output #0: loss =
0.689023 (* 1 = 0.689023 loss)
I0907 17:49:19.840370 2018 sgd_solver.cpp:106] Iteration 540, lr = 0.001
I0907 17:49:20.117620 2018 solver.cpp:229] Iteration 560, loss = 0.691357
I0907 17:49:20.117674 2018 solver.cpp:245] Train net output #0: loss =
0.691357 (* 1 = 0.691357 loss)
I0907 17:49:20.117691 2018 sgd_solver.cpp:106] Iteration 560, lr = 0.001
I0907 17:49:20.395112 2018 solver.cpp:229] Iteration 580, loss = 0.686913
I0907 17:49:20.395167 2018 solver.cpp:245] Train net output #0: loss =
0.686913 (* 1 = 0.686913 loss)
I0907 17:49:20.395184 2018 sgd_solver.cpp:106] Iteration 580, lr = 0.001
I0907 17:49:20.672291 2018 solver.cpp:229] Iteration 600, loss = 0.688692
I0907 17:49:20.672343 2018 solver.cpp:245] Train net output #0: loss =
0.688692 (* 1 = 0.688692 loss)
I0907 17:49:20.672363 2018 sgd_solver.cpp:106] Iteration 600, lr = 0.001
I0907 17:49:20.949833 2018 solver.cpp:229] Iteration 620, loss = 0.690929
I0907 17:49:20.949885 2018 solver.cpp:245] Train net output #0: loss =
0.690929 (* 1 = 0.690929 loss)
I0907 17:49:20.949900 2018 sgd_solver.cpp:106] Iteration 620, lr = 0.001
I0907 17:49:21.227128 2018 solver.cpp:229] Iteration 640, loss = 0.696944
I0907 17:49:21.227190 2018 solver.cpp:245] Train net output #0: loss =

```

```

0.696944 (* 1 = 0.696944 loss)
I0907 17:49:21.227205 2018 sgd_solver.cpp:106] Iteration 640, lr = 0.001
I0907 17:49:21.504741 2018 solver.cpp:229] Iteration 660, loss = 0.690737
I0907 17:49:21.504804 2018 solver.cpp:245] Train net output #0: loss =
0.690737 (* 1 = 0.690737 loss)
I0907 17:49:21.504820 2018 sgd_solver.cpp:106] Iteration 660, lr = 0.001
I0907 17:49:21.781639 2018 solver.cpp:229] Iteration 680, loss = 0.696515
I0907 17:49:21.781697 2018 solver.cpp:245] Train net output #0: loss =
0.696515 (* 1 = 0.696515 loss)
I0907 17:49:21.781713 2018 sgd_solver.cpp:106] Iteration 680, lr = 0.001
I0907 17:49:22.058897 2018 solver.cpp:229] Iteration 700, loss = 0.696229
I0907 17:49:22.058950 2018 solver.cpp:245] Train net output #0: loss =
0.696229 (* 1 = 0.696229 loss)
I0907 17:49:22.058965 2018 sgd_solver.cpp:106] Iteration 700, lr = 0.001
I0907 17:49:22.336318 2018 solver.cpp:229] Iteration 720, loss = 0.692999
I0907 17:49:22.336371 2018 solver.cpp:245] Train net output #0: loss =
0.692999 (* 1 = 0.692999 loss)
I0907 17:49:22.336386 2018 sgd_solver.cpp:106] Iteration 720, lr = 0.001
I0907 17:49:22.613425 2018 solver.cpp:229] Iteration 740, loss = 0.690792
I0907 17:49:22.613484 2018 solver.cpp:245] Train net output #0: loss =
0.690792 (* 1 = 0.690792 loss)
I0907 17:49:22.613499 2018 sgd_solver.cpp:106] Iteration 740, lr = 0.001
I0907 17:49:22.738293 2018 solver.cpp:456] Snapshotting to binary proto file
checkpoints/snapshot_iter_750.caffemodel
I0907 17:49:22.748927 2018 sgd_solver.cpp:273] Snapshotting solver state to
binary proto file checkpoints/snapshot_iter_750.solverstate
I0907 17:49:22.750164 2018 solver.cpp:338] Iteration 750, Testing net (#0)
I0907 17:49:23.240234 2018 solver.cpp:406] Test net output #0: accuracy
= 0.5054
I0907 17:49:23.240293 2018 solver.cpp:406] Test net output #1: loss = 0.
693245 (* 1 = 0.693245 loss)
I0907 17:49:23.240309 2018 solver.cpp:323] Optimization Done.
I0907 17:49:23.240315 2018 caffe.cpp:216] Optimization Done.

```

After 750 training iterations the training accuracy is about 50%. Not impressive!

We can achieve much higher accuracy and a lower loss against both the training and validation datasets for this task. It seems our network is underfitting the data.

## Refinements

Many network configurations such as Alexnet, GoogLeNet and VGG are significantly larger than the three layer architecture used above and have proven to be very accurate at classifying the ImageNet images. Now, we are going to increase the complexity of this network to improve the accuracy.

There are many knobs that one can turn in choosing a neural network architecture. For example, add layers, increased the number of learned weights, change the learning rate or introduce a more complex policy to modify the learning rate as training progresses. We will experiment with some of these modifications to see the effect on classification accuracy.

## Refinement #1

We will increase the number of outputs in the convolutional layers to 64 for layer 1 and 128 for layer 2. Also add a third convolutional layer that is identical to the second one. Just make sure to have the right layer outputs feeding into this new layer and the final fully-connected layer.

This time it will take about 90 seconds to train due to the increased network size.

```
In [9]: %%bash  
TOOLS=/home/ubuntu/caffe/build/tools  
#Train your modified network configuration  
$TOOLS/caffe train -gpu 0 -solver src/solver2.prototxt
```



```
I0907 17:49:56.266579 2026 upgrade_proto.cpp:990] Attempting to upgrade input file specified using deprecated 'solver_type' field (enum)': src/solver2.prototxt
I0907 17:49:56.266804 2026 upgrade_proto.cpp:997] Successfully upgraded file specified using deprecated 'solver_type' field (enum) to 'type' field (string).
W0907 17:49:56.266813 2026 upgrade_proto.cpp:999] Note that future Caffe releases will only support 'type' field (string) for a solver's type.
I0907 17:49:56.266942 2026 caffe.cpp:185] Using GPUs 0
I0907 17:49:56.397658 2026 solver.cpp:48] Initializing solver from parameters:
test_iter: 100
test_interval: 250
base_lr: 0.01
display: 20
max_iter: 750
lr_policy: "step"
gamma: 0.1
momentum: 0.9
weight_decay: 0.0005
stepsize: 500
snapshot: 250
snapshot_prefix: "checkpoints/snapshot"
solver_mode: GPU
device_id: 0
random_seed: 1234
net: "src/train_val2.answer.prototxt"
type: "SGD"
I0907 17:49:56.397824 2026 solver.cpp:91] Creating training net from net file: src/train_val2.answer.prototxt
I0907 17:49:56.398375 2026 net.cpp:322] The NetState phase (0) differed from the phase (1) specified by a rule in layer data
I0907 17:49:56.398409 2026 net.cpp:322] The NetState phase (0) differed from the phase (1) specified by a rule in layer accuracy
I0907 17:49:56.398548 2026 net.cpp:49] Initializing net from parameters:
state {
    phase: TRAIN
}
layer {
    name: "data"
    type: "Data"
    top: "data"
    top: "label"
    include {
        phase: TRAIN
    }
    transform_param {
        mirror: true
        crop_size: 32
        mean_file: "mean.binaryproto"
    }
    data_param {
        source: "train_lmdb"
        batch_size: 100
        backend: LMDB
    }
}
```

```
layer {
    name: "conv1"
    type: "Convolution"
    bottom: "data"
    top: "conv1"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 64
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.0001
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "pool1"
    type: "Pooling"
    bottom: "conv1"
    top: "pool1"
    pooling_param {
        pool: MAX
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "relu1"
    type: "ReLU"
    bottom: "pool1"
    top: "pool1"
}
layer {
    name: "conv2"
    type: "Convolution"
    bottom: "pool1"
    top: "conv2"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 128
        pad: 2
        kernel_size: 5
    }
}
```

```
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "relu2"
    type: "ReLU"
    bottom: "conv2"
    top: "conv2"
}
layer {
    name: "pool2"
    type: "Pooling"
    bottom: "conv2"
    top: "pool2"
    pooling_param {
        pool: AVE
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "conv3"
    type: "Convolution"
    bottom: "pool2"
    top: "conv3"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 128
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "relu3"
    type: "ReLU"
    bottom: "conv3"
    top: "conv3"
```

```

}
layer {
    name: "pool3"
    type: "Pooling"
    bottom: "conv3"
    top: "pool3"
    pooling_param {
        pool: AVE
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "ip1"
    type: "InnerProduct"
    bottom: "pool3"
    top: "ip1"
    param {
        lr_mult: 1
        decay_mult: 250
    }
    param {
        lr_mult: 2
        decay_mult: 0
    }
    inner_product_param {
        num_output: 2
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "loss"
    type: "SoftmaxWithLoss"
    bottom: "ip1"
    bottom: "label"
    top: "loss"
}
I0907 17:49:56.398650 2026 layer_factory.hpp:77] Creating layer data
I0907 17:49:56.398857 2026 net.cpp:106] Creating Layer data
I0907 17:49:56.398879 2026 net.cpp:411] data -> data
I0907 17:49:56.398928 2026 net.cpp:411] data -> label
I0907 17:49:56.398957 2026 data_transformer.cpp:25] Loading mean file from:
mean.binaryproto
I0907 17:49:56.399595 2029 db_lmdb.cpp:38] Opened lmdb train_lmdb
I0907 17:49:56.411839 2026 data_layer.cpp:41] output data size: 100,3,32,32
I0907 17:49:56.414516 2026 net.cpp:150] Setting up data
I0907 17:49:56.414572 2026 net.cpp:157] Top shape: 100 3 32 32 (307200)
I0907 17:49:56.414580 2026 net.cpp:157] Top shape: 100 (100)
I0907 17:49:56.414584 2026 net.cpp:165] Memory required for data: 1229200
I0907 17:49:56.414595 2026 layer_factory.hpp:77] Creating layer conv1
I0907 17:49:56.414625 2026 net.cpp:106] Creating Layer conv1

```

```
I0907 17:49:56.414636 2026 net.cpp:454] conv1 <- data
I0907 17:49:56.414657 2026 net.cpp:411] conv1 -> conv1
I0907 17:49:56.531949 2026 net.cpp:150] Setting up conv1
I0907 17:49:56.531997 2026 net.cpp:157] Top shape: 100 64 32 32 (6553600)
I0907 17:49:56.532004 2026 net.cpp:165] Memory required for data: 27443600
I0907 17:49:56.532033 2026 layer_factory.hpp:77] Creating layer pool1
I0907 17:49:56.532063 2026 net.cpp:106] Creating Layer pool1
I0907 17:49:56.532078 2026 net.cpp:454] pool1 <- conv1
I0907 17:49:56.532088 2026 net.cpp:411] pool1 -> pool1
I0907 17:49:56.532163 2026 net.cpp:150] Setting up pool1
I0907 17:49:56.532181 2026 net.cpp:157] Top shape: 100 64 16 16 (1638400)
I0907 17:49:56.532186 2026 net.cpp:165] Memory required for data: 33997200
I0907 17:49:56.532194 2026 layer_factory.hpp:77] Creating layer relu1
I0907 17:49:56.532203 2026 net.cpp:106] Creating Layer relu1
I0907 17:49:56.532213 2026 net.cpp:454] relu1 <- pool1
I0907 17:49:56.532222 2026 net.cpp:397] relu1 -> pool1 (in-place)
I0907 17:49:56.532369 2026 net.cpp:150] Setting up relu1
I0907 17:49:56.532388 2026 net.cpp:157] Top shape: 100 64 16 16 (1638400)
I0907 17:49:56.532394 2026 net.cpp:165] Memory required for data: 40550800
I0907 17:49:56.532399 2026 layer_factory.hpp:77] Creating layer conv2
I0907 17:49:56.532420 2026 net.cpp:106] Creating Layer conv2
I0907 17:49:56.532433 2026 net.cpp:454] conv2 <- pool1
I0907 17:49:56.532441 2026 net.cpp:411] conv2 -> conv2
I0907 17:49:56.540318 2026 net.cpp:150] Setting up conv2
I0907 17:49:56.540345 2026 net.cpp:157] Top shape: 100 128 16 16 (3276800)
I0907 17:49:56.540357 2026 net.cpp:165] Memory required for data: 53658000
I0907 17:49:56.540370 2026 layer_factory.hpp:77] Creating layer relu2
I0907 17:49:56.540385 2026 net.cpp:106] Creating Layer relu2
I0907 17:49:56.540390 2026 net.cpp:454] relu2 <- conv2
I0907 17:49:56.540407 2026 net.cpp:397] relu2 -> conv2 (in-place)
I0907 17:49:56.540650 2026 net.cpp:150] Setting up relu2
I0907 17:49:56.540671 2026 net.cpp:157] Top shape: 100 128 16 16 (3276800)
I0907 17:49:56.540676 2026 net.cpp:165] Memory required for data: 66765200
I0907 17:49:56.540689 2026 layer_factory.hpp:77] Creating layer pool2
I0907 17:49:56.540701 2026 net.cpp:106] Creating Layer pool2
I0907 17:49:56.540712 2026 net.cpp:454] pool2 <- conv2
I0907 17:49:56.540720 2026 net.cpp:411] pool2 -> pool2
I0907 17:49:56.540904 2026 net.cpp:150] Setting up pool2
I0907 17:49:56.540925 2026 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:49:56.540930 2026 net.cpp:165] Memory required for data: 70042000
I0907 17:49:56.540940 2026 layer_factory.hpp:77] Creating layer conv3
I0907 17:49:56.540956 2026 net.cpp:106] Creating Layer conv3
I0907 17:49:56.540966 2026 net.cpp:454] conv3 <- pool2
I0907 17:49:56.540978 2026 net.cpp:411] conv3 -> conv3
I0907 17:49:56.555501 2026 net.cpp:150] Setting up conv3
I0907 17:49:56.555526 2026 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:49:56.555532 2026 net.cpp:165] Memory required for data: 73318800
I0907 17:49:56.555548 2026 layer_factory.hpp:77] Creating layer relu3
I0907 17:49:56.555560 2026 net.cpp:106] Creating Layer relu3
I0907 17:49:56.555567 2026 net.cpp:454] relu3 <- conv3
I0907 17:49:56.555577 2026 net.cpp:397] relu3 -> conv3 (in-place)
I0907 17:49:56.555745 2026 net.cpp:150] Setting up relu3
I0907 17:49:56.555764 2026 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:49:56.555770 2026 net.cpp:165] Memory required for data: 76595600
I0907 17:49:56.555778 2026 layer_factory.hpp:77] Creating layer pool3
I0907 17:49:56.555788 2026 net.cpp:106] Creating Layer pool3
I0907 17:49:56.555799 2026 net.cpp:454] pool3 <- conv3
```

```
I0907 17:49:56.555843 2026 net.cpp:411] pool3 -> pool3
I0907 17:49:56.556097 2026 net.cpp:150] Setting up pool3
I0907 17:49:56.556119 2026 net.cpp:157] Top shape: 100 128 4 4 (204800)
I0907 17:49:56.556125 2026 net.cpp:165] Memory required for data: 77414800
I0907 17:49:56.556131 2026 layer_factory.hpp:77] Creating layer ip1
I0907 17:49:56.556148 2026 net.cpp:106] Creating Layer ip1
I0907 17:49:56.556159 2026 net.cpp:454] ip1 <- pool3
I0907 17:49:56.556174 2026 net.cpp:411] ip1 -> ip1
I0907 17:49:56.556427 2026 net.cpp:150] Setting up ip1
I0907 17:49:56.556444 2026 net.cpp:157] Top shape: 100 2 (200)
I0907 17:49:56.556449 2026 net.cpp:165] Memory required for data: 77415600
I0907 17:49:56.556462 2026 layer_factory.hpp:77] Creating layer loss
I0907 17:49:56.556475 2026 net.cpp:106] Creating Layer loss
I0907 17:49:56.556486 2026 net.cpp:454] loss <- ip1
I0907 17:49:56.556493 2026 net.cpp:454] loss <- label
I0907 17:49:56.556509 2026 net.cpp:411] loss -> loss
I0907 17:49:56.556526 2026 layer_factory.hpp:77] Creating layer loss
I0907 17:49:56.556879 2026 net.cpp:150] Setting up loss
I0907 17:49:56.556900 2026 net.cpp:157] Top shape: (1)
I0907 17:49:56.556905 2026 net.cpp:160] with loss weight 1
I0907 17:49:56.556936 2026 net.cpp:165] Memory required for data: 77415604
I0907 17:49:56.556942 2026 net.cpp:226] loss needs backward computation.
I0907 17:49:56.556948 2026 net.cpp:226] ip1 needs backward computation.
I0907 17:49:56.556954 2026 net.cpp:226] pool3 needs backward computation.
I0907 17:49:56.556958 2026 net.cpp:226] relu3 needs backward computation.
I0907 17:49:56.556964 2026 net.cpp:226] conv3 needs backward computation.
I0907 17:49:56.556969 2026 net.cpp:226] pool2 needs backward computation.
I0907 17:49:56.556979 2026 net.cpp:226] relu2 needs backward computation.
I0907 17:49:56.556984 2026 net.cpp:226] conv2 needs backward computation.
I0907 17:49:56.556988 2026 net.cpp:226] relu1 needs backward computation.
I0907 17:49:56.556993 2026 net.cpp:226] pool1 needs backward computation.
I0907 17:49:56.556998 2026 net.cpp:226] conv1 needs backward computation.
I0907 17:49:56.557004 2026 net.cpp:228] data does not need backward computation.

I0907 17:49:56.557009 2026 net.cpp:270] This network produces output loss
I0907 17:49:56.557026 2026 net.cpp:283] Network initialization done.
I0907 17:49:56.557523 2026 solver.cpp:181] Creating test net (#0) specified
    by net file: src/train_val2.answer.prototxt
I0907 17:49:56.557579 2026 net.cpp:322] The NetState phase (1) differed from
    the phase (0) specified by a rule in layer data
I0907 17:49:56.557720 2026 net.cpp:49] Initializing net from parameters:
state {
    phase: TEST
}
layer {
    name: "data"
    type: "Data"
    top: "data"
    top: "label"
    include {
        phase: TEST
    }
    transform_param {
        crop_size: 32
        mean_file: "mean.binaryproto"
    }
    data_param {
```

```
source: "val_lmdb"
batch_size: 100
backend: LMDB
}
}
layer {
    name: "conv1"
    type: "Convolution"
    bottom: "data"
    top: "conv1"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 64
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.0001
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "pool1"
    type: "Pooling"
    bottom: "conv1"
    top: "pool1"
    pooling_param {
        pool: MAX
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "relu1"
    type: "ReLU"
    bottom: "pool1"
    top: "pool1"
}
layer {
    name: "conv2"
    type: "Convolution"
    bottom: "pool1"
    top: "conv2"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
}
```

```
        }
    convolution_param {
        num_output: 128
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}

layer {
    name: "relu2"
    type: "ReLU"
    bottom: "conv2"
    top: "conv2"
}
layer {
    name: "pool2"
    type: "Pooling"
    bottom: "conv2"
    top: "pool2"
    pooling_param {
        pool: AVE
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "conv3"
    type: "Convolution"
    bottom: "pool2"
    top: "conv3"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 128
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}
```

```
layer {
    name: "relu3"
    type: "ReLU"
    bottom: "conv3"
    top: "conv3"
}
layer {
    name: "pool3"
    type: "Pooling"
    bottom: "conv3"
    top: "pool3"
    pooling_param {
        pool: AVE
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "ip1"
    type: "InnerProduct"
    bottom: "pool3"
    top: "ip1"
    param {
        lr_mult: 1
        decay_mult: 250
    }
    param {
        lr_mult: 2
        decay_mult: 0
    }
    inner_product_param {
        num_output: 2
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "loss"
    type: "SoftmaxWithLoss"
    bottom: "ip1"
    bottom: "label"
    top: "loss"
}
layer {
    name: "accuracy"
    type: "Accuracy"
    bottom: "ip1"
    bottom: "label"
    top: "accuracy"
    include {
        phase: TEST
    }
}
```

```
}
```

I0907 17:49:56.557842 2026 layer\_factory.hpp:77] Creating layer data  
I0907 17:49:56.557986 2026 net.cpp:106] Creating Layer data  
I0907 17:49:56.558002 2026 net.cpp:411] data -> data  
I0907 17:49:56.558014 2026 net.cpp:411] data -> label  
I0907 17:49:56.558037 2026 data\_transformer.cpp:25] Loading mean file from:  
mean.binaryproto  
I0907 17:49:56.558761 2031 db\_lmdb.cpp:38] Opened lmdb val\_lmdb  
I0907 17:49:56.558881 2026 data\_layer.cpp:41] output data size: 100,3,32,32  
I0907 17:49:56.561842 2026 net.cpp:150] Setting up data  
I0907 17:49:56.561864 2026 net.cpp:157] Top shape: 100 3 32 32 (307200)  
I0907 17:49:56.561872 2026 net.cpp:157] Top shape: 100 (100)  
I0907 17:49:56.561877 2026 net.cpp:165] Memory required for data: 1229200  
I0907 17:49:56.561882 2026 layer\_factory.hpp:77] Creating layer label\_data\_1  
\_split  
I0907 17:49:56.561890 2026 net.cpp:106] Creating Layer label\_data\_1\_split  
I0907 17:49:56.561897 2026 net.cpp:454] label\_data\_1\_split <- label  
I0907 17:49:56.561904 2026 net.cpp:411] label\_data\_1\_split -> label\_data\_1\_s  
plit\_0  
I0907 17:49:56.561918 2026 net.cpp:411] label\_data\_1\_split -> label\_data\_1\_s  
plit\_1  
I0907 17:49:56.561998 2026 net.cpp:150] Setting up label\_data\_1\_split  
I0907 17:49:56.562016 2026 net.cpp:157] Top shape: 100 (100)  
I0907 17:49:56.562022 2026 net.cpp:157] Top shape: 100 (100)  
I0907 17:49:56.562026 2026 net.cpp:165] Memory required for data: 1230000  
I0907 17:49:56.562031 2026 layer\_factory.hpp:77] Creating layer conv1  
I0907 17:49:56.562055 2026 net.cpp:106] Creating Layer conv1  
I0907 17:49:56.562062 2026 net.cpp:454] conv1 <- data  
I0907 17:49:56.562073 2026 net.cpp:411] conv1 -> conv1  
I0907 17:49:56.563552 2026 net.cpp:150] Setting up conv1  
I0907 17:49:56.563575 2026 net.cpp:157] Top shape: 100 64 32 32 (6553600)  
I0907 17:49:56.563580 2026 net.cpp:165] Memory required for data: 27444400  
I0907 17:49:56.563621 2026 layer\_factory.hpp:77] Creating layer pool1  
I0907 17:49:56.563632 2026 net.cpp:106] Creating Layer pool1  
I0907 17:49:56.563639 2026 net.cpp:454] pool1 <- conv1  
I0907 17:49:56.563649 2026 net.cpp:411] pool1 -> pool1  
I0907 17:49:56.563697 2026 net.cpp:150] Setting up pool1  
I0907 17:49:56.563714 2026 net.cpp:157] Top shape: 100 64 16 16 (1638400)  
I0907 17:49:56.563719 2026 net.cpp:165] Memory required for data: 33998000  
I0907 17:49:56.563724 2026 layer\_factory.hpp:77] Creating layer relu1  
I0907 17:49:56.563735 2026 net.cpp:106] Creating Layer relu1  
I0907 17:49:56.563742 2026 net.cpp:454] relu1 <- pool1  
I0907 17:49:56.563750 2026 net.cpp:397] relu1 -> pool1 (in-place)  
I0907 17:49:56.564000 2026 net.cpp:150] Setting up relu1  
I0907 17:49:56.564020 2026 net.cpp:157] Top shape: 100 64 16 16 (1638400)  
I0907 17:49:56.564028 2026 net.cpp:165] Memory required for data: 40551600  
I0907 17:49:56.564034 2026 layer\_factory.hpp:77] Creating layer conv2  
I0907 17:49:56.564045 2026 net.cpp:106] Creating Layer conv2  
I0907 17:49:56.564050 2026 net.cpp:454] conv2 <- pool1  
I0907 17:49:56.564061 2026 net.cpp:411] conv2 -> conv2  
I0907 17:49:56.571995 2026 net.cpp:150] Setting up conv2  
I0907 17:49:56.572021 2026 net.cpp:157] Top shape: 100 128 16 16 (3276800)  
I0907 17:49:56.572026 2026 net.cpp:165] Memory required for data: 53658800  
I0907 17:49:56.572038 2026 layer\_factory.hpp:77] Creating layer relu2  
I0907 17:49:56.572053 2026 net.cpp:106] Creating Layer relu2  
I0907 17:49:56.572064 2026 net.cpp:454] relu2 <- conv2  
I0907 17:49:56.572075 2026 net.cpp:397] relu2 -> conv2 (in-place)

```
I0907 17:49:56.572238 2026 net.cpp:150] Setting up relu2
I0907 17:49:56.572257 2026 net.cpp:157] Top shape: 100 128 16 16 (3276800)
I0907 17:49:56.572263 2026 net.cpp:165] Memory required for data: 66766000
I0907 17:49:56.572274 2026 layer_factory.hpp:77] Creating layer pool2
I0907 17:49:56.572283 2026 net.cpp:106] Creating Layer pool2
I0907 17:49:56.572295 2026 net.cpp:454] pool2 <- conv2
I0907 17:49:56.572304 2026 net.cpp:411] pool2 -> pool2
I0907 17:49:56.572567 2026 net.cpp:150] Setting up pool2
I0907 17:49:56.572588 2026 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:49:56.572593 2026 net.cpp:165] Memory required for data: 70042800
I0907 17:49:56.572600 2026 layer_factory.hpp:77] Creating layer conv3
I0907 17:49:56.572618 2026 net.cpp:106] Creating Layer conv3
I0907 17:49:56.572629 2026 net.cpp:454] conv3 <- pool2
I0907 17:49:56.572638 2026 net.cpp:411] conv3 -> conv3
I0907 17:49:56.588001 2026 net.cpp:150] Setting up conv3
I0907 17:49:56.588032 2026 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:49:56.588037 2026 net.cpp:165] Memory required for data: 73319600
I0907 17:49:56.588058 2026 layer_factory.hpp:77] Creating layer relu3
I0907 17:49:56.588073 2026 net.cpp:106] Creating Layer relu3
I0907 17:49:56.588079 2026 net.cpp:454] relu3 <- conv3
I0907 17:49:56.588091 2026 net.cpp:397] relu3 -> conv3 (in-place)
I0907 17:49:56.588265 2026 net.cpp:150] Setting up relu3
I0907 17:49:56.588284 2026 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:49:56.588290 2026 net.cpp:165] Memory required for data: 76596400
I0907 17:49:56.588295 2026 layer_factory.hpp:77] Creating layer pool3
I0907 17:49:56.588307 2026 net.cpp:106] Creating Layer pool3
I0907 17:49:56.588317 2026 net.cpp:454] pool3 <- conv3
I0907 17:49:56.588328 2026 net.cpp:411] pool3 -> pool3
I0907 17:49:56.588639 2026 net.cpp:150] Setting up pool3
I0907 17:49:56.588657 2026 net.cpp:157] Top shape: 100 128 4 4 (204800)
I0907 17:49:56.588663 2026 net.cpp:165] Memory required for data: 77415600
I0907 17:49:56.588668 2026 layer_factory.hpp:77] Creating layer ip1
I0907 17:49:56.588683 2026 net.cpp:106] Creating Layer ip1
I0907 17:49:56.588691 2026 net.cpp:454] ip1 <- pool3
I0907 17:49:56.588701 2026 net.cpp:411] ip1 -> ip1
I0907 17:49:56.588956 2026 net.cpp:150] Setting up ip1
I0907 17:49:56.588976 2026 net.cpp:157] Top shape: 100 2 (200)
I0907 17:49:56.588981 2026 net.cpp:165] Memory required for data: 77416400
I0907 17:49:56.588999 2026 layer_factory.hpp:77] Creating layer ip1_ip1_0_sp
lit
I0907 17:49:56.589013 2026 net.cpp:106] Creating Layer ip1_ip1_0_split
I0907 17:49:56.589020 2026 net.cpp:454] ip1_ip1_0_split <- ip1
I0907 17:49:56.589027 2026 net.cpp:411] ip1_ip1_0_split -> ip1_ip1_0_split_0
I0907 17:49:56.589035 2026 net.cpp:411] ip1_ip1_0_split -> ip1_ip1_0_split_1
I0907 17:49:56.589088 2026 net.cpp:150] Setting up ip1_ip1_0_split
I0907 17:49:56.589103 2026 net.cpp:157] Top shape: 100 2 (200)
I0907 17:49:56.589110 2026 net.cpp:157] Top shape: 100 2 (200)
I0907 17:49:56.589117 2026 net.cpp:165] Memory required for data: 77418000
I0907 17:49:56.589121 2026 layer_factory.hpp:77] Creating layer loss
I0907 17:49:56.589129 2026 net.cpp:106] Creating Layer loss
I0907 17:49:56.589139 2026 net.cpp:454] loss <- ip1_ip1_0_split_0
I0907 17:49:56.589145 2026 net.cpp:454] loss <- label_data_1_split_0
I0907 17:49:56.589156 2026 net.cpp:411] loss -> loss
I0907 17:49:56.589169 2026 layer_factory.hpp:77] Creating layer loss
I0907 17:49:56.589413 2026 net.cpp:150] Setting up loss
I0907 17:49:56.589433 2026 net.cpp:157] Top shape: (1)
I0907 17:49:56.589462 2026 net.cpp:160] with loss weight 1
```

```
I0907 17:49:56.589478 2026 net.cpp:165] Memory required for data: 77418004
I0907 17:49:56.589483 2026 layer_factory.hpp:77] Creating layer accuracy
I0907 17:49:56.589498 2026 net.cpp:106] Creating Layer accuracy
I0907 17:49:56.589504 2026 net.cpp:454] accuracy <- ip1_ip1_0_split_1
I0907 17:49:56.589510 2026 net.cpp:454] accuracy <- label_data_1_split_1
I0907 17:49:56.589521 2026 net.cpp:411] accuracy -> accuracy
I0907 17:49:56.589540 2026 net.cpp:150] Setting up accuracy
I0907 17:49:56.589552 2026 net.cpp:157] Top shape: (1)
I0907 17:49:56.589557 2026 net.cpp:165] Memory required for data: 77418008
I0907 17:49:56.589562 2026 net.cpp:228] accuracy does not need backward computation.
I0907 17:49:56.589570 2026 net.cpp:226] loss needs backward computation.
I0907 17:49:56.589576 2026 net.cpp:226] ip1_ip1_0_split needs backward computation.
I0907 17:49:56.589579 2026 net.cpp:226] ip1 needs backward computation.
I0907 17:49:56.589586 2026 net.cpp:226] pool3 needs backward computation.
I0907 17:49:56.589591 2026 net.cpp:226] relu3 needs backward computation.
I0907 17:49:56.589594 2026 net.cpp:226] conv3 needs backward computation.
I0907 17:49:56.589601 2026 net.cpp:226] pool2 needs backward computation.
I0907 17:49:56.589606 2026 net.cpp:226] relu2 needs backward computation.
I0907 17:49:56.589610 2026 net.cpp:226] conv2 needs backward computation.
I0907 17:49:56.589617 2026 net.cpp:226] relu1 needs backward computation.
I0907 17:49:56.589622 2026 net.cpp:226] pool1 needs backward computation.
I0907 17:49:56.589625 2026 net.cpp:226] conv1 needs backward computation.
I0907 17:49:56.589633 2026 net.cpp:228] label_data_1_split does not need backward computation.
I0907 17:49:56.589638 2026 net.cpp:228] data does not need backward computation.
I0907 17:49:56.589643 2026 net.cpp:270] This network produces output accuracy
I0907 17:49:56.589648 2026 net.cpp:270] This network produces output loss
I0907 17:49:56.589663 2026 net.cpp:283] Network initialization done.
I0907 17:49:56.589745 2026 solver.cpp:60] Solver scaffolding done.
I0907 17:49:56.590039 2026 caffe.cpp:213] Starting Optimization
I0907 17:49:56.590062 2026 solver.cpp:280] Solving
I0907 17:49:56.590067 2026 solver.cpp:281] Learning Rate Policy: step
I0907 17:49:56.590569 2026 solver.cpp:338] Iteration 0, Testing net (#0)
I0907 17:49:58.996317 2026 solver.cpp:406] Test net output #0: accuracy = 0.4942
I0907 17:49:58.996376 2026 solver.cpp:406] Test net output #1: loss = 0.693116 (* 1 = 0.693116 loss)
I0907 17:49:59.022239 2026 solver.cpp:229] Iteration 0, loss = 0.692804
I0907 17:49:59.022272 2026 solver.cpp:245] Train net output #0: loss = 0.692804 (* 1 = 0.692804 loss)
I0907 17:49:59.022292 2026 sgd_solver.cpp:106] Iteration 0, lr = 0.01
I0907 17:50:01.016228 2026 solver.cpp:229] Iteration 20, loss = 0.681038
I0907 17:50:01.016283 2026 solver.cpp:245] Train net output #0: loss = 0.681038 (* 1 = 0.681038 loss)
I0907 17:50:01.016294 2026 sgd_solver.cpp:106] Iteration 20, lr = 0.01
I0907 17:50:03.008795 2026 solver.cpp:229] Iteration 40, loss = 0.67226
I0907 17:50:03.008852 2026 solver.cpp:245] Train net output #0: loss = 0.67226 (* 1 = 0.67226 loss)
I0907 17:50:03.008862 2026 sgd_solver.cpp:106] Iteration 40, lr = 0.01
I0907 17:50:05.001580 2026 solver.cpp:229] Iteration 60, loss = 0.673935
I0907 17:50:05.001646 2026 solver.cpp:245] Train net output #0: loss = 0.673935 (* 1 = 0.673935 loss)
I0907 17:50:05.001655 2026 sgd_solver.cpp:106] Iteration 60, lr = 0.01
```

```

I0907 17:50:06.994484 2026 solver.cpp:229] Iteration 80, loss = 0.692611
I0907 17:50:06.994542 2026 solver.cpp:245]      Train net output #0: loss =
  0.692611 (* 1 = 0.692611 loss)
I0907 17:50:06.994554 2026 sgd_solver.cpp:106] Iteration 80, lr = 0.01
I0907 17:50:08.988448 2026 solver.cpp:229] Iteration 100, loss = 0.674448
I0907 17:50:08.988512 2026 solver.cpp:245]      Train net output #0: loss =
  0.674448 (* 1 = 0.674448 loss)
I0907 17:50:08.988523 2026 sgd_solver.cpp:106] Iteration 100, lr = 0.01
I0907 17:50:10.983147 2026 solver.cpp:229] Iteration 120, loss = 0.683734
I0907 17:50:10.983203 2026 solver.cpp:245]      Train net output #0: loss =
  0.683734 (* 1 = 0.683734 loss)
I0907 17:50:10.983213 2026 sgd_solver.cpp:106] Iteration 120, lr = 0.01
I0907 17:50:12.977910 2026 solver.cpp:229] Iteration 140, loss = 0.63254
I0907 17:50:12.977974 2026 solver.cpp:245]      Train net output #0: loss =
  0.63254 (* 1 = 0.63254 loss)
I0907 17:50:12.977984 2026 sgd_solver.cpp:106] Iteration 140, lr = 0.01
I0907 17:50:14.970306 2026 solver.cpp:229] Iteration 160, loss = 0.624842
I0907 17:50:14.970360 2026 solver.cpp:245]      Train net output #0: loss =
  0.624842 (* 1 = 0.624842 loss)
I0907 17:50:14.970371 2026 sgd_solver.cpp:106] Iteration 160, lr = 0.01
I0907 17:50:16.963953 2026 solver.cpp:229] Iteration 180, loss = 0.597834
I0907 17:50:16.964012 2026 solver.cpp:245]      Train net output #0: loss =
  0.597834 (* 1 = 0.597834 loss)
I0907 17:50:16.964023 2026 sgd_solver.cpp:106] Iteration 180, lr = 0.01
I0907 17:50:18.958662 2026 solver.cpp:229] Iteration 200, loss = 0.59141
I0907 17:50:18.958729 2026 solver.cpp:245]      Train net output #0: loss =
  0.59141 (* 1 = 0.59141 loss)
I0907 17:50:18.958745 2026 sgd_solver.cpp:106] Iteration 200, lr = 0.01
I0907 17:50:20.951462 2026 solver.cpp:229] Iteration 220, loss = 0.66259
I0907 17:50:20.951515 2026 solver.cpp:245]      Train net output #0: loss =
  0.66259 (* 1 = 0.66259 loss)
I0907 17:50:20.951526 2026 sgd_solver.cpp:106] Iteration 220, lr = 0.01
I0907 17:50:22.946105 2026 solver.cpp:229] Iteration 240, loss = 0.6401
I0907 17:50:22.946171 2026 solver.cpp:245]      Train net output #0: loss =
  0.6401 (* 1 = 0.6401 loss)
I0907 17:50:22.946182 2026 sgd_solver.cpp:106] Iteration 240, lr = 0.01
I0907 17:50:23.843701 2026 solver.cpp:456] Snapshotting to binary proto file
  checkpoints/snapshot_iter_250.caffemodel
I0907 17:50:23.933241 2026 sgd_solver.cpp:273] Snapshotting solver state to
  binary proto file checkpoints/snapshot_iter_250.solverstate
I0907 17:50:23.938874 2026 solver.cpp:338] Iteration 250, Testing net (#0)
I0907 17:50:26.336319 2026 solver.cpp:406]      Test net output #0: accuracy
  = 0.6662
I0907 17:50:26.336613 2026 solver.cpp:406]      Test net output #1: loss = 0.
  623816 (* 1 = 0.623816 loss)
I0907 17:50:27.358484 2026 solver.cpp:229] Iteration 260, loss = 0.69391
I0907 17:50:27.358542 2026 solver.cpp:245]      Train net output #0: loss =
  0.69391 (* 1 = 0.69391 loss)
I0907 17:50:27.358553 2026 sgd_solver.cpp:106] Iteration 260, lr = 0.01
I0907 17:50:29.352551 2026 solver.cpp:229] Iteration 280, loss = 0.637233
I0907 17:50:29.352607 2026 solver.cpp:245]      Train net output #0: loss =
  0.637233 (* 1 = 0.637233 loss)
I0907 17:50:29.352617 2026 sgd_solver.cpp:106] Iteration 280, lr = 0.01
I0907 17:50:31.346138 2026 solver.cpp:229] Iteration 300, loss = 0.537481
I0907 17:50:31.346195 2026 solver.cpp:245]      Train net output #0: loss =
  0.537481 (* 1 = 0.537481 loss)
I0907 17:50:31.346206 2026 sgd_solver.cpp:106] Iteration 300, lr = 0.01

```

```
I0907 17:50:33.340416 2026 solver.cpp:229] Iteration 320, loss = 0.630526
I0907 17:50:33.340476 2026 solver.cpp:245] Train net output #0: loss =
0.630526 (* 1 = 0.630526 loss)
I0907 17:50:33.340489 2026 sgd_solver.cpp:106] Iteration 320, lr = 0.01
I0907 17:50:35.330626 2026 solver.cpp:229] Iteration 340, loss = 0.591059
I0907 17:50:35.330682 2026 solver.cpp:245] Train net output #0: loss =
0.591059 (* 1 = 0.591059 loss)
I0907 17:50:35.330693 2026 sgd_solver.cpp:106] Iteration 340, lr = 0.01
I0907 17:50:37.322973 2026 solver.cpp:229] Iteration 360, loss = 0.593058
I0907 17:50:37.323031 2026 solver.cpp:245] Train net output #0: loss =
0.593058 (* 1 = 0.593058 loss)
I0907 17:50:37.323050 2026 sgd_solver.cpp:106] Iteration 360, lr = 0.01
I0907 17:50:39.315917 2026 solver.cpp:229] Iteration 380, loss = 0.61844
I0907 17:50:39.315976 2026 solver.cpp:245] Train net output #0: loss =
0.61844 (* 1 = 0.61844 loss)
I0907 17:50:39.315986 2026 sgd_solver.cpp:106] Iteration 380, lr = 0.01
I0907 17:50:41.309324 2026 solver.cpp:229] Iteration 400, loss = 0.690648
I0907 17:50:41.309386 2026 solver.cpp:245] Train net output #0: loss =
0.690648 (* 1 = 0.690648 loss)
I0907 17:50:41.309396 2026 sgd_solver.cpp:106] Iteration 400, lr = 0.01
I0907 17:50:43.301913 2026 solver.cpp:229] Iteration 420, loss = 0.62693
I0907 17:50:43.301976 2026 solver.cpp:245] Train net output #0: loss =
0.62693 (* 1 = 0.62693 loss)
I0907 17:50:43.301990 2026 sgd_solver.cpp:106] Iteration 420, lr = 0.01
I0907 17:50:45.294443 2026 solver.cpp:229] Iteration 440, loss = 0.600443
I0907 17:50:45.294498 2026 solver.cpp:245] Train net output #0: loss =
0.600443 (* 1 = 0.600443 loss)
I0907 17:50:45.294508 2026 sgd_solver.cpp:106] Iteration 440, lr = 0.01
I0907 17:50:47.288780 2026 solver.cpp:229] Iteration 460, loss = 0.700381
I0907 17:50:47.288833 2026 solver.cpp:245] Train net output #0: loss =
0.700381 (* 1 = 0.700381 loss)
I0907 17:50:47.288843 2026 sgd_solver.cpp:106] Iteration 460, lr = 0.01
I0907 17:50:49.281090 2026 solver.cpp:229] Iteration 480, loss = 0.555135
I0907 17:50:49.281150 2026 solver.cpp:245] Train net output #0: loss =
0.555135 (* 1 = 0.555135 loss)
I0907 17:50:49.281162 2026 sgd_solver.cpp:106] Iteration 480, lr = 0.01
I0907 17:50:51.175283 2026 solver.cpp:456] Snapshotting to binary proto file
checkpoints/snapshot_iter_500.caffemodel
I0907 17:50:51.262532 2026 sgd_solver.cpp:273] Snapshotting solver state to
binary proto file checkpoints/snapshot_iter_500.solverstate
I0907 17:50:51.267627 2026 solver.cpp:338] Iteration 500, Testing net (#0)
I0907 17:50:53.665534 2026 solver.cpp:406] Test net output #0: accuracy
= 0.661
I0907 17:50:53.665586 2026 solver.cpp:406] Test net output #1: loss = 0.
628322 (* 1 = 0.628322 loss)
I0907 17:50:53.689847 2026 solver.cpp:229] Iteration 500, loss = 0.584232
I0907 17:50:53.689877 2026 solver.cpp:245] Train net output #0: loss =
0.584232 (* 1 = 0.584232 loss)
I0907 17:50:53.689888 2026 sgd_solver.cpp:106] Iteration 500, lr = 0.001
I0907 17:50:55.685356 2026 solver.cpp:229] Iteration 520, loss = 0.621351
I0907 17:50:55.685413 2026 solver.cpp:245] Train net output #0: loss =
0.621351 (* 1 = 0.621351 loss)
I0907 17:50:55.685462 2026 sgd_solver.cpp:106] Iteration 520, lr = 0.001
I0907 17:50:57.679355 2026 solver.cpp:229] Iteration 540, loss = 0.587776
I0907 17:50:57.679518 2026 solver.cpp:245] Train net output #0: loss =
0.587776 (* 1 = 0.587776 loss)
I0907 17:50:57.679530 2026 sgd_solver.cpp:106] Iteration 540, lr = 0.001
```

```
I0907 17:50:59.672078 2026 solver.cpp:229] Iteration 560, loss = 0.547018
I0907 17:50:59.672145 2026 solver.cpp:245]      Train net output #0: loss =
0.547018 (* 1 = 0.547018 loss)
I0907 17:50:59.672155 2026 sgd_solver.cpp:106] Iteration 560, lr = 0.001
I0907 17:51:01.665143 2026 solver.cpp:229] Iteration 580, loss = 0.612046
I0907 17:51:01.665199 2026 solver.cpp:245]      Train net output #0: loss =
0.612046 (* 1 = 0.612046 loss)
I0907 17:51:01.665210 2026 sgd_solver.cpp:106] Iteration 580, lr = 0.001
I0907 17:51:03.659373 2026 solver.cpp:229] Iteration 600, loss = 0.651342
I0907 17:51:03.659430 2026 solver.cpp:245]      Train net output #0: loss =
0.651342 (* 1 = 0.651342 loss)
I0907 17:51:03.659441 2026 sgd_solver.cpp:106] Iteration 600, lr = 0.001
I0907 17:51:05.654263 2026 solver.cpp:229] Iteration 620, loss = 0.5456
I0907 17:51:05.654319 2026 solver.cpp:245]      Train net output #0: loss =
0.5456 (* 1 = 0.5456 loss)
I0907 17:51:05.654328 2026 sgd_solver.cpp:106] Iteration 620, lr = 0.001
I0907 17:51:07.647342 2026 solver.cpp:229] Iteration 640, loss = 0.616289
I0907 17:51:07.647402 2026 solver.cpp:245]      Train net output #0: loss =
0.616289 (* 1 = 0.616289 loss)
I0907 17:51:07.647413 2026 sgd_solver.cpp:106] Iteration 640, lr = 0.001
I0907 17:51:09.638375 2026 solver.cpp:229] Iteration 660, loss = 0.570817
I0907 17:51:09.638432 2026 solver.cpp:245]      Train net output #0: loss =
0.570817 (* 1 = 0.570817 loss)
I0907 17:51:09.638442 2026 sgd_solver.cpp:106] Iteration 660, lr = 0.001
I0907 17:51:11.631057 2026 solver.cpp:229] Iteration 680, loss = 0.550111
I0907 17:51:11.631124 2026 solver.cpp:245]      Train net output #0: loss =
0.550111 (* 1 = 0.550111 loss)
I0907 17:51:11.631135 2026 sgd_solver.cpp:106] Iteration 680, lr = 0.001
I0907 17:51:13.623296 2026 solver.cpp:229] Iteration 700, loss = 0.54983
I0907 17:51:13.623353 2026 solver.cpp:245]      Train net output #0: loss =
0.54983 (* 1 = 0.54983 loss)
I0907 17:51:13.623364 2026 sgd_solver.cpp:106] Iteration 700, lr = 0.001
I0907 17:51:15.616344 2026 solver.cpp:229] Iteration 720, loss = 0.565444
I0907 17:51:15.616401 2026 solver.cpp:245]      Train net output #0: loss =
0.565444 (* 1 = 0.565444 loss)
I0907 17:51:15.616411 2026 sgd_solver.cpp:106] Iteration 720, lr = 0.001
I0907 17:51:17.612784 2026 solver.cpp:229] Iteration 740, loss = 0.522849
I0907 17:51:17.612841 2026 solver.cpp:245]      Train net output #0: loss =
0.522849 (* 1 = 0.522849 loss)
I0907 17:51:17.612851 2026 sgd_solver.cpp:106] Iteration 740, lr = 0.001
I0907 17:51:18.510428 2026 solver.cpp:456] Snapshotting to binary proto file
checkpoints/snapshot_iter_750.caffemodel
I0907 17:51:18.597537 2026 sgd_solver.cpp:273] Snapshotting solver state to
binary proto file checkpoints/snapshot_iter_750.solverstate
I0907 17:51:18.602501 2026 solver.cpp:338] Iteration 750, Testing net (#0)
I0907 17:51:21.004096 2026 solver.cpp:406]      Test net output #0: accuracy
= 0.7128
I0907 17:51:21.004165 2026 solver.cpp:406]      Test net output #1: loss = 0.
571076 (* 1 = 0.571076 loss)
I0907 17:51:21.004176 2026 solver.cpp:323] Optimization Done.
I0907 17:51:21.004181 2026 caffe.cpp:216] Optimization Done.
```

We noticed significant performance improvements. The accuracy is now about 71% and we are underfitting less.

## Refinement #2

Let's make one more modification to the network before training a final time. In the last modification we increased the number of neurons in our convolutional layers. Another way to increase the number of trainable parameters in our network is to make it deeper by adding more layers. This time we add a new fully-connected layer with 500 outputs - call it ip2. This layer should come after the third pooling layer, pool3, but **before** the existing fully-connected layer ip1. In Caffe, fully-connected layers are implemented using the inner product layer construct. After the new fully-connected (inner product) layer you also need a ReLU activation layer and a dropout layer. The dropout layer will prevent the network from *overfitting*, i.e. getting really good at classifying the training data but not able to classify the validation data.

Again, this model will take slightly longer to train due to the increased size. It's about 3 minutes and 20 seconds.

```
In [10]: %%bash  
TOOLS=/home/ubuntu/caffe/build/tools  
#Train your modified network configuration  
$TOOLS/caffe train -gpu 0 -solver src/solver3.prototxt
```



```
I0907 17:52:18.155535 2034 upgrade_proto.cpp:990] Attempting to upgrade input file specified using deprecated 'solver_type' field (enum)': src/solver3.prototxt
I0907 17:52:18.155802 2034 upgrade_proto.cpp:997] Successfully upgraded file specified using deprecated 'solver_type' field (enum) to 'type' field (string).
W0907 17:52:18.155814 2034 upgrade_proto.cpp:999] Note that future Caffe releases will only support 'type' field (string) for a solver's type.
I0907 17:52:18.155977 2034 caffe.cpp:185] Using GPUs 0
I0907 17:52:18.290024 2034 solver.cpp:48] Initializing solver from parameters:
test_iter: 100
test_interval: 250
base_lr: 0.01
display: 20
max_iter: 750
lr_policy: "step"
gamma: 0.1
momentum: 0.9
stepsize: 500
snapshot: 750
snapshot_prefix: "checkpoints/snapshot"
solver_mode: GPU
device_id: 0
random_seed: 1234
net: "src/train_val3.answer.prototxt"
type: "SGD"
I0907 17:52:18.290211 2034 solver.cpp:91] Creating training net from net file: src/train_val3.answer.prototxt
I0907 17:52:18.290791 2034 net.cpp:322] The NetState phase (0) differed from the phase (1) specified by a rule in layer data
I0907 17:52:18.290827 2034 net.cpp:322] The NetState phase (0) differed from the phase (1) specified by a rule in layer accuracy
I0907 17:52:18.290977 2034 net.cpp:49] Initializing net from parameters:
state {
    phase: TRAIN
}
layer {
    name: "data"
    type: "Data"
    top: "data"
    top: "label"
    include {
        phase: TRAIN
    }
    transform_param {
        mirror: true
        crop_size: 32
        mean_file: "mean.binaryproto"
    }
    data_param {
        source: "train_lmdb"
        batch_size: 100
        backend: LMDB
    }
}
layer {
```

```
name: "conv1"
type: "Convolution"
bottom: "data"
top: "conv1"
param {
    lr_mult: 1
}
param {
    lr_mult: 2
}
convolution_param {
    num_output: 64
    pad: 2
    kernel_size: 5
    stride: 1
    weight_filler {
        type: "gaussian"
        std: 0.0001
    }
    bias_filler {
        type: "constant"
    }
}
layer {
    name: "pool1"
    type: "Pooling"
    bottom: "conv1"
    top: "pool1"
    pooling_param {
        pool: MAX
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "relu1"
    type: "ReLU"
    bottom: "pool1"
    top: "pool1"
}
layer {
    name: "conv2"
    type: "Convolution"
    bottom: "pool1"
    top: "conv2"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 128
        pad: 2
        kernel_size: 5
        stride: 1
    }
}
```

```
weight_filler {
    type: "gaussian"
    std: 0.01
}
bias_filler {
    type: "constant"
}
}
layer {
    name: "relu2"
    type: "ReLU"
    bottom: "conv2"
    top: "conv2"
}
layer {
    name: "pool2"
    type: "Pooling"
    bottom: "conv2"
    top: "pool2"
    pooling_param {
        pool: AVE
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "conv3"
    type: "Convolution"
    bottom: "pool2"
    top: "conv3"
    convolution_param {
        num_output: 128
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "relu3"
    type: "ReLU"
    bottom: "conv3"
    top: "conv3"
}
layer {
    name: "pool3"
    type: "Pooling"
    bottom: "conv3"
    top: "pool3"
    pooling_param {
```

```
        pool: AVE
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "ip2"
    type: "InnerProduct"
    bottom: "pool3"
    top: "ip2"
    param {
        lr_mult: 1
        decay_mult: 1
    }
    param {
        lr_mult: 2
        decay_mult: 0
    }
    inner_product_param {
        num_output: 500
        weight_filler {
            type: "gaussian"
            std: 0.005
        }
        bias_filler {
            type: "constant"
            value: 0.1
        }
    }
}
layer {
    name: "reluip2"
    type: "ReLU"
    bottom: "ip2"
    top: "ip2"
}
layer {
    name: "dropip2"
    type: "Dropout"
    bottom: "ip2"
    top: "ip2"
    dropout_param {
        dropout_ratio: 0.5
    }
}
layer {
    name: "ip1"
    type: "InnerProduct"
    bottom: "ip2"
    top: "ip1"
    param {
        lr_mult: 1
        decay_mult: 250
    }
    param {
        lr_mult: 2
        decay_mult: 0
    }
}
```

```
        }
        inner_product_param {
            num_output: 2
            weight_filler {
                type: "gaussian"
                std: 0.01
            }
            bias_filler {
                type: "constant"
            }
        }
    }
layer {
    name: "loss"
    type: "SoftmaxWithLoss"
    bottom: "ip1"
    bottom: "label"
    top: "loss"
}
I0907 17:52:18.291097 2034 layer_factory.hpp:77] Creating layer data
I0907 17:52:18.291326 2034 net.cpp:106] Creating Layer data
I0907 17:52:18.291352 2034 net.cpp:411] data -> data
I0907 17:52:18.291399 2034 net.cpp:411] data -> label
I0907 17:52:18.291429 2034 data_transformer.cpp:25] Loading mean file from:
mean.binaryproto
I0907 17:52:18.292052 2037 db_lmdb.cpp:38] Opened lmdb train_lmdb
I0907 17:52:18.303288 2034 data_layer.cpp:41] output data size: 100,3,32,32
I0907 17:52:18.306143 2034 net.cpp:150] Setting up data
I0907 17:52:18.306174 2034 net.cpp:157] Top shape: 100 3 32 32 (307200)
I0907 17:52:18.306185 2034 net.cpp:157] Top shape: 100 (100)
I0907 17:52:18.306190 2034 net.cpp:165] Memory required for data: 1229200
I0907 17:52:18.306201 2034 layer_factory.hpp:77] Creating layer conv1
I0907 17:52:18.306229 2034 net.cpp:106] Creating Layer conv1
I0907 17:52:18.306246 2034 net.cpp:454] conv1 <- data
I0907 17:52:18.306267 2034 net.cpp:411] conv1 -> conv1
I0907 17:52:18.424154 2034 net.cpp:150] Setting up conv1
I0907 17:52:18.424204 2034 net.cpp:157] Top shape: 100 64 32 32 (6553600)
I0907 17:52:18.424212 2034 net.cpp:165] Memory required for data: 27443600
I0907 17:52:18.424240 2034 layer_factory.hpp:77] Creating layer pool1
I0907 17:52:18.424263 2034 net.cpp:106] Creating Layer pool1
I0907 17:52:18.424278 2034 net.cpp:454] pool1 <- conv1
I0907 17:52:18.424288 2034 net.cpp:411] pool1 -> pool1
I0907 17:52:18.424353 2034 net.cpp:150] Setting up pool1
I0907 17:52:18.424371 2034 net.cpp:157] Top shape: 100 64 16 16 (1638400)
I0907 17:52:18.424376 2034 net.cpp:165] Memory required for data: 33997200
I0907 17:52:18.424381 2034 layer_factory.hpp:77] Creating layer relu1
I0907 17:52:18.424393 2034 net.cpp:106] Creating Layer relu1
I0907 17:52:18.424401 2034 net.cpp:454] relu1 <- pool1
I0907 17:52:18.424408 2034 net.cpp:397] relu1 -> pool1 (in-place)
I0907 17:52:18.424556 2034 net.cpp:150] Setting up relu1
I0907 17:52:18.424574 2034 net.cpp:157] Top shape: 100 64 16 16 (1638400)
I0907 17:52:18.424579 2034 net.cpp:165] Memory required for data: 40550800
I0907 17:52:18.424590 2034 layer_factory.hpp:77] Creating layer conv2
I0907 17:52:18.424609 2034 net.cpp:106] Creating Layer conv2
I0907 17:52:18.424619 2034 net.cpp:454] conv2 <- pool1
I0907 17:52:18.424628 2034 net.cpp:411] conv2 -> conv2
I0907 17:52:18.432421 2034 net.cpp:150] Setting up conv2
```

```
I0907 17:52:18.432446 2034 net.cpp:157] Top shape: 100 128 16 16 (3276800)
I0907 17:52:18.432452 2034 net.cpp:165] Memory required for data: 53658000
I0907 17:52:18.432476 2034 layer_factory.hpp:77] Creating layer relu2
I0907 17:52:18.432487 2034 net.cpp:106] Creating Layer relu2
I0907 17:52:18.432497 2034 net.cpp:454] relu2 <- conv2
I0907 17:52:18.432503 2034 net.cpp:397] relu2 -> conv2 (in-place)
I0907 17:52:18.432729 2034 net.cpp:150] Setting up relu2
I0907 17:52:18.432750 2034 net.cpp:157] Top shape: 100 128 16 16 (3276800)
I0907 17:52:18.432756 2034 net.cpp:165] Memory required for data: 66765200
I0907 17:52:18.432772 2034 layer_factory.hpp:77] Creating layer pool2
I0907 17:52:18.432781 2034 net.cpp:106] Creating Layer pool2
I0907 17:52:18.432792 2034 net.cpp:454] pool2 <- conv2
I0907 17:52:18.432801 2034 net.cpp:411] pool2 -> pool2
I0907 17:52:18.432976 2034 net.cpp:150] Setting up pool2
I0907 17:52:18.432997 2034 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:52:18.433006 2034 net.cpp:165] Memory required for data: 70042000
I0907 17:52:18.433012 2034 layer_factory.hpp:77] Creating layer conv3
I0907 17:52:18.433025 2034 net.cpp:106] Creating Layer conv3
I0907 17:52:18.433037 2034 net.cpp:454] conv3 <- pool2
I0907 17:52:18.433045 2034 net.cpp:411] conv3 -> conv3
I0907 17:52:18.447526 2034 net.cpp:150] Setting up conv3
I0907 17:52:18.447549 2034 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:52:18.447556 2034 net.cpp:165] Memory required for data: 73318800
I0907 17:52:18.447571 2034 layer_factory.hpp:77] Creating layer relu3
I0907 17:52:18.447597 2034 net.cpp:106] Creating Layer relu3
I0907 17:52:18.447609 2034 net.cpp:454] relu3 <- conv3
I0907 17:52:18.447655 2034 net.cpp:397] relu3 -> conv3 (in-place)
I0907 17:52:18.447814 2034 net.cpp:150] Setting up relu3
I0907 17:52:18.447834 2034 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:52:18.447839 2034 net.cpp:165] Memory required for data: 76595600
I0907 17:52:18.447845 2034 layer_factory.hpp:77] Creating layer pool3
I0907 17:52:18.447854 2034 net.cpp:106] Creating Layer pool3
I0907 17:52:18.447865 2034 net.cpp:454] pool3 <- conv3
I0907 17:52:18.447875 2034 net.cpp:411] pool3 -> pool3
I0907 17:52:18.448151 2034 net.cpp:150] Setting up pool3
I0907 17:52:18.448171 2034 net.cpp:157] Top shape: 100 128 4 4 (204800)
I0907 17:52:18.448177 2034 net.cpp:165] Memory required for data: 77414800
I0907 17:52:18.448185 2034 layer_factory.hpp:77] Creating layer ip2
I0907 17:52:18.448199 2034 net.cpp:106] Creating Layer ip2
I0907 17:52:18.448210 2034 net.cpp:454] ip2 <- pool3
I0907 17:52:18.448221 2034 net.cpp:411] ip2 -> ip2
I0907 17:52:18.482381 2034 net.cpp:150] Setting up ip2
I0907 17:52:18.482406 2034 net.cpp:157] Top shape: 100 500 (50000)
I0907 17:52:18.482412 2034 net.cpp:165] Memory required for data: 77614800
I0907 17:52:18.482421 2034 layer_factory.hpp:77] Creating layer reluip2
I0907 17:52:18.482439 2034 net.cpp:106] Creating Layer reluip2
I0907 17:52:18.482455 2034 net.cpp:454] reluip2 <- ip2
I0907 17:52:18.482463 2034 net.cpp:397] reluip2 -> ip2 (in-place)
I0907 17:52:18.482728 2034 net.cpp:150] Setting up reluip2
I0907 17:52:18.482748 2034 net.cpp:157] Top shape: 100 500 (50000)
I0907 17:52:18.482754 2034 net.cpp:165] Memory required for data: 77814800
I0907 17:52:18.482761 2034 layer_factory.hpp:77] Creating layer dropip2
I0907 17:52:18.482774 2034 net.cpp:106] Creating Layer dropip2
I0907 17:52:18.482784 2034 net.cpp:454] dropip2 <- ip2
I0907 17:52:18.482795 2034 net.cpp:397] dropip2 -> ip2 (in-place)
I0907 17:52:18.482841 2034 net.cpp:150] Setting up dropip2
I0907 17:52:18.482857 2034 net.cpp:157] Top shape: 100 500 (50000)
```

```
I0907 17:52:18.482866 2034 net.cpp:165] Memory required for data: 78014800
I0907 17:52:18.482872 2034 layer_factory.hpp:77] Creating layer ip1
I0907 17:52:18.482882 2034 net.cpp:106] Creating Layer ip1
I0907 17:52:18.482892 2034 net.cpp:454] ip1 <- ip2
I0907 17:52:18.482900 2034 net.cpp:411] ip1 -> ip1
I0907 17:52:18.483049 2034 net.cpp:150] Setting up ip1
I0907 17:52:18.483067 2034 net.cpp:157] Top shape: 100 2 (200)
I0907 17:52:18.483072 2034 net.cpp:165] Memory required for data: 78015600
I0907 17:52:18.483088 2034 layer_factory.hpp:77] Creating layer loss
I0907 17:52:18.483114 2034 net.cpp:106] Creating Layer loss
I0907 17:52:18.483125 2034 net.cpp:454] loss <- ip1
I0907 17:52:18.483131 2034 net.cpp:454] loss <- label
I0907 17:52:18.483144 2034 net.cpp:411] loss -> loss
I0907 17:52:18.483165 2034 layer_factory.hpp:77] Creating layer loss
I0907 17:52:18.483433 2034 net.cpp:150] Setting up loss
I0907 17:52:18.483451 2034 net.cpp:157] Top shape: (1)
I0907 17:52:18.483456 2034 net.cpp:160]      with loss weight 1
I0907 17:52:18.483489 2034 net.cpp:165] Memory required for data: 78015604
I0907 17:52:18.483496 2034 net.cpp:226] loss needs backward computation.
I0907 17:52:18.483502 2034 net.cpp:226] ip1 needs backward computation.
I0907 17:52:18.483508 2034 net.cpp:226] dropip2 needs backward computation.
I0907 17:52:18.483512 2034 net.cpp:226] reluip2 needs backward computation.
I0907 17:52:18.483518 2034 net.cpp:226] ip2 needs backward computation.
I0907 17:52:18.483523 2034 net.cpp:226] pool3 needs backward computation.
I0907 17:52:18.483528 2034 net.cpp:226] relu3 needs backward computation.
I0907 17:52:18.483533 2034 net.cpp:226] conv3 needs backward computation.
I0907 17:52:18.483539 2034 net.cpp:226] pool2 needs backward computation.
I0907 17:52:18.483543 2034 net.cpp:226] relu2 needs backward computation.
I0907 17:52:18.483549 2034 net.cpp:226] conv2 needs backward computation.
I0907 17:52:18.483553 2034 net.cpp:226] relu1 needs backward computation.
I0907 17:52:18.483559 2034 net.cpp:226] pool1 needs backward computation.
I0907 17:52:18.483595 2034 net.cpp:226] conv1 needs backward computation.
I0907 17:52:18.483610 2034 net.cpp:228] data does not need backward computation.

I0907 17:52:18.483616 2034 net.cpp:270] This network produces output loss
I0907 17:52:18.483631 2034 net.cpp:283] Network initialization done.
I0907 17:52:18.484194 2034 solver.cpp:181] Creating test net (#0) specified
    by net file: src/train_val3.answer.prototxt
I0907 17:52:18.484259 2034 net.cpp:322] The NetState phase (1) differed from
    the phase (0) specified by a rule in layer data
I0907 17:52:18.484424 2034 net.cpp:49] Initializing net from parameters:
state {
    phase: TEST
}
layer {
    name: "data"
    type: "Data"
    top: "data"
    top: "label"
    include {
        phase: TEST
    }
    transform_param {
        crop_size: 32
        mean_file: "mean.binaryproto"
    }
    data_param {
```

```
source: "val_lmdb"
batch_size: 100
backend: LMDB
}
}
layer {
    name: "conv1"
    type: "Convolution"
    bottom: "data"
    top: "conv1"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 64
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.0001
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "pool1"
    type: "Pooling"
    bottom: "conv1"
    top: "pool1"
    pooling_param {
        pool: MAX
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "relu1"
    type: "ReLU"
    bottom: "pool1"
    top: "pool1"
}
layer {
    name: "conv2"
    type: "Convolution"
    bottom: "pool1"
    top: "conv2"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
}
```

```
        }
    convolution_param {
        num_output: 128
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}

layer {
    name: "relu2"
    type: "ReLU"
    bottom: "conv2"
    top: "conv2"
}

layer {
    name: "pool2"
    type: "Pooling"
    bottom: "conv2"
    top: "pool2"
    pooling_param {
        pool: AVE
        kernel_size: 3
        stride: 2
    }
}

layer {
    name: "conv3"
    type: "Convolution"
    bottom: "pool2"
    top: "conv3"
    convolution_param {
        num_output: 128
        pad: 2
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}

layer {
    name: "relu3"
    type: "ReLU"
    bottom: "conv3"
    top: "conv3"
}
```

```
layer {
    name: "pool3"
    type: "Pooling"
    bottom: "conv3"
    top: "pool3"
    pooling_param {
        pool: AVE
        kernel_size: 3
        stride: 2
    }
}
layer {
    name: "ip2"
    type: "InnerProduct"
    bottom: "pool3"
    top: "ip2"
    param {
        lr_mult: 1
        decay_mult: 1
    }
    param {
        lr_mult: 2
        decay_mult: 0
    }
    inner_product_param {
        num_output: 500
        weight_filler {
            type: "gaussian"
            std: 0.005
        }
        bias_filler {
            type: "constant"
            value: 0.1
        }
    }
}
layer {
    name: "reluip2"
    type: "ReLU"
    bottom: "ip2"
    top: "ip2"
}
layer {
    name: "dropip2"
    type: "Dropout"
    bottom: "ip2"
    top: "ip2"
    dropout_param {
        dropout_ratio: 0.5
    }
}
layer {
    name: "ip1"
    type: "InnerProduct"
    bottom: "ip2"
    top: "ip1"
    param {
```

```

        lr_mult: 1
        decay_mult: 250
    }
    param {
        lr_mult: 2
        decay_mult: 0
    }
    inner_product_param {
        num_output: 2
        weight_filler {
            type: "gaussian"
            std: 0.01
        }
        bias_filler {
            type: "constant"
        }
    }
}
layer {
    name: "loss"
    type: "SoftmaxWithLoss"
    bottom: "ip1"
    bottom: "label"
    top: "loss"
}
layer {
    name: "accuracy"
    type: "Accuracy"
    bottom: "ip1"
    bottom: "label"
    top: "accuracy"
    include {
        phase: TEST
    }
}
I0907 17:52:18.484539 2034 layer_factory.hpp:77] Creating layer data
I0907 17:52:18.484683 2034 net.cpp:106] Creating Layer data
I0907 17:52:18.484700 2034 net.cpp:411] data -> data
I0907 17:52:18.484714 2034 net.cpp:411] data -> label
I0907 17:52:18.484736 2034 data_transformer.cpp:25] Loading mean file from:
mean.binaryproto
I0907 17:52:18.485389 2039 db_lmdb.cpp:38] Opened lmdb val_lmdb
I0907 17:52:18.485507 2034 data_layer.cpp:41] output data size: 100,3,32,32
I0907 17:52:18.488396 2034 net.cpp:150] Setting up data
I0907 17:52:18.488420 2034 net.cpp:157] Top shape: 100 3 32 32 (307200)
I0907 17:52:18.488427 2034 net.cpp:157] Top shape: 100 (100)
I0907 17:52:18.488432 2034 net.cpp:165] Memory required for data: 1229200
I0907 17:52:18.488437 2034 layer_factory.hpp:77] Creating layer label_data_1_split
I0907 17:52:18.488450 2034 net.cpp:106] Creating Layer label_data_1_split
I0907 17:52:18.488456 2034 net.cpp:454] label_data_1_split <- label
I0907 17:52:18.488467 2034 net.cpp:411] label_data_1_split -> label_data_1_s
plit_0
I0907 17:52:18.488477 2034 net.cpp:411] label_data_1_split -> label_data_1_s
plit_1
I0907 17:52:18.488661 2034 net.cpp:150] Setting up label_data_1_split
I0907 17:52:18.488688 2034 net.cpp:157] Top shape: 100 (100)

```

```
I0907 17:52:18.488695 2034 net.cpp:157] Top shape: 100 (100)
I0907 17:52:18.488699 2034 net.cpp:165] Memory required for data: 1230000
I0907 17:52:18.488704 2034 layer_factory.hpp:77] Creating layer conv1
I0907 17:52:18.488719 2034 net.cpp:106] Creating Layer conv1
I0907 17:52:18.488731 2034 net.cpp:454] conv1 <- data
I0907 17:52:18.488744 2034 net.cpp:411] conv1 -> conv1
I0907 17:52:18.490237 2034 net.cpp:150] Setting up conv1
I0907 17:52:18.490260 2034 net.cpp:157] Top shape: 100 64 32 32 (6553600)
I0907 17:52:18.490267 2034 net.cpp:165] Memory required for data: 27444400
I0907 17:52:18.490279 2034 layer_factory.hpp:77] Creating layer pool1
I0907 17:52:18.490293 2034 net.cpp:106] Creating Layer pool1
I0907 17:52:18.490299 2034 net.cpp:454] pool1 <- conv1
I0907 17:52:18.490308 2034 net.cpp:411] pool1 -> pool1
I0907 17:52:18.490356 2034 net.cpp:150] Setting up pool1
I0907 17:52:18.490373 2034 net.cpp:157] Top shape: 100 64 16 16 (1638400)
I0907 17:52:18.490378 2034 net.cpp:165] Memory required for data: 33998000
I0907 17:52:18.490384 2034 layer_factory.hpp:77] Creating layer relu1
I0907 17:52:18.490396 2034 net.cpp:106] Creating Layer relu1
I0907 17:52:18.490402 2034 net.cpp:454] relu1 <- pool1
I0907 17:52:18.490409 2034 net.cpp:397] relu1 -> pool1 (in-place)
I0907 17:52:18.490564 2034 net.cpp:150] Setting up relu1
I0907 17:52:18.490583 2034 net.cpp:157] Top shape: 100 64 16 16 (1638400)
I0907 17:52:18.490589 2034 net.cpp:165] Memory required for data: 40551600
I0907 17:52:18.490594 2034 layer_factory.hpp:77] Creating layer conv2
I0907 17:52:18.490608 2034 net.cpp:106] Creating Layer conv2
I0907 17:52:18.490617 2034 net.cpp:454] conv2 <- pool1
I0907 17:52:18.490628 2034 net.cpp:411] conv2 -> conv2
I0907 17:52:18.498924 2034 net.cpp:150] Setting up conv2
I0907 17:52:18.498950 2034 net.cpp:157] Top shape: 100 128 16 16 (3276800)
I0907 17:52:18.498955 2034 net.cpp:165] Memory required for data: 53658800
I0907 17:52:18.498970 2034 layer_factory.hpp:77] Creating layer relu2
I0907 17:52:18.498991 2034 net.cpp:106] Creating Layer relu2
I0907 17:52:18.499001 2034 net.cpp:454] relu2 <- conv2
I0907 17:52:18.499008 2034 net.cpp:397] relu2 -> conv2 (in-place)
I0907 17:52:18.499266 2034 net.cpp:150] Setting up relu2
I0907 17:52:18.499287 2034 net.cpp:157] Top shape: 100 128 16 16 (3276800)
I0907 17:52:18.499294 2034 net.cpp:165] Memory required for data: 66766000
I0907 17:52:18.499300 2034 layer_factory.hpp:77] Creating layer pool2
I0907 17:52:18.499311 2034 net.cpp:106] Creating Layer pool2
I0907 17:52:18.499321 2034 net.cpp:454] pool2 <- conv2
I0907 17:52:18.499330 2034 net.cpp:411] pool2 -> pool2
I0907 17:52:18.499505 2034 net.cpp:150] Setting up pool2
I0907 17:52:18.499524 2034 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:52:18.499534 2034 net.cpp:165] Memory required for data: 70042800
I0907 17:52:18.499544 2034 layer_factory.hpp:77] Creating layer conv3
I0907 17:52:18.499562 2034 net.cpp:106] Creating Layer conv3
I0907 17:52:18.499573 2034 net.cpp:454] conv3 <- pool2
I0907 17:52:18.499608 2034 net.cpp:411] conv3 -> conv3
I0907 17:52:18.514439 2034 net.cpp:150] Setting up conv3
I0907 17:52:18.514464 2034 net.cpp:157] Top shape: 100 128 8 8 (819200)
I0907 17:52:18.514470 2034 net.cpp:165] Memory required for data: 73319600
I0907 17:52:18.514513 2034 layer_factory.hpp:77] Creating layer relu3
I0907 17:52:18.514526 2034 net.cpp:106] Creating Layer relu3
I0907 17:52:18.514539 2034 net.cpp:454] relu3 <- conv3
I0907 17:52:18.514546 2034 net.cpp:397] relu3 -> conv3 (in-place)
I0907 17:52:18.514808 2034 net.cpp:150] Setting up relu3
I0907 17:52:18.514829 2034 net.cpp:157] Top shape: 100 128 8 8 (819200)
```

```
I0907 17:52:18.514835 2034 net.cpp:165] Memory required for data: 76596400
I0907 17:52:18.514842 2034 layer_factory.hpp:77] Creating layer pool3
I0907 17:52:18.514852 2034 net.cpp:106] Creating Layer pool3
I0907 17:52:18.514861 2034 net.cpp:454] pool3 <- conv3
I0907 17:52:18.514873 2034 net.cpp:411] pool3 -> pool3
I0907 17:52:18.515054 2034 net.cpp:150] Setting up pool3
I0907 17:52:18.515074 2034 net.cpp:157] Top shape: 100 128 4 4 (204800)
I0907 17:52:18.515079 2034 net.cpp:165] Memory required for data: 77415600
I0907 17:52:18.515085 2034 layer_factory.hpp:77] Creating layer ip2
I0907 17:52:18.515095 2034 net.cpp:106] Creating Layer ip2
I0907 17:52:18.515106 2034 net.cpp:454] ip2 <- pool3
I0907 17:52:18.515117 2034 net.cpp:411] ip2 -> ip2
I0907 17:52:18.549113 2034 net.cpp:150] Setting up ip2
I0907 17:52:18.549139 2034 net.cpp:157] Top shape: 100 500 (50000)
I0907 17:52:18.549144 2034 net.cpp:165] Memory required for data: 77615600
I0907 17:52:18.549154 2034 layer_factory.hpp:77] Creating layer reluip2
I0907 17:52:18.549168 2034 net.cpp:106] Creating Layer reluip2
I0907 17:52:18.549175 2034 net.cpp:454] reluip2 <- ip2
I0907 17:52:18.549183 2034 net.cpp:397] reluip2 -> ip2 (in-place)
I0907 17:52:18.549437 2034 net.cpp:150] Setting up reluip2
I0907 17:52:18.549458 2034 net.cpp:157] Top shape: 100 500 (50000)
I0907 17:52:18.549463 2034 net.cpp:165] Memory required for data: 77815600
I0907 17:52:18.549470 2034 layer_factory.hpp:77] Creating layer dropip2
I0907 17:52:18.549479 2034 net.cpp:106] Creating Layer dropip2
I0907 17:52:18.549489 2034 net.cpp:454] dropip2 <- ip2
I0907 17:52:18.549500 2034 net.cpp:397] dropip2 -> ip2 (in-place)
I0907 17:52:18.549540 2034 net.cpp:150] Setting up dropip2
I0907 17:52:18.549558 2034 net.cpp:157] Top shape: 100 500 (50000)
I0907 17:52:18.549563 2034 net.cpp:165] Memory required for data: 78015600
I0907 17:52:18.549576 2034 layer_factory.hpp:77] Creating layer ip1
I0907 17:52:18.549587 2034 net.cpp:106] Creating Layer ip1
I0907 17:52:18.549597 2034 net.cpp:454] ip1 <- ip2
I0907 17:52:18.549608 2034 net.cpp:411] ip1 -> ip1
I0907 17:52:18.549767 2034 net.cpp:150] Setting up ip1
I0907 17:52:18.549784 2034 net.cpp:157] Top shape: 100 2 (200)
I0907 17:52:18.549789 2034 net.cpp:165] Memory required for data: 78016400
I0907 17:52:18.549803 2034 layer_factory.hpp:77] Creating layer ip1_ip1_0_sp
lit
I0907 17:52:18.549818 2034 net.cpp:106] Creating Layer ip1_ip1_0_split
I0907 17:52:18.549824 2034 net.cpp:454] ip1_ip1_0_split <- ip1
I0907 17:52:18.549839 2034 net.cpp:411] ip1_ip1_0_split -> ip1_ip1_0_split_0
I0907 17:52:18.549849 2034 net.cpp:411] ip1_ip1_0_split -> ip1_ip1_0_split_1
I0907 17:52:18.549897 2034 net.cpp:150] Setting up ip1_ip1_0_split
I0907 17:52:18.549916 2034 net.cpp:157] Top shape: 100 2 (200)
I0907 17:52:18.549921 2034 net.cpp:157] Top shape: 100 2 (200)
I0907 17:52:18.549927 2034 net.cpp:165] Memory required for data: 78018000
I0907 17:52:18.549932 2034 layer_factory.hpp:77] Creating layer loss
I0907 17:52:18.549947 2034 net.cpp:106] Creating Layer loss
I0907 17:52:18.549953 2034 net.cpp:454] loss <- ip1_ip1_0_split_0
I0907 17:52:18.549960 2034 net.cpp:454] loss <- label_data_1_split_0
I0907 17:52:18.549968 2034 net.cpp:411] loss -> loss
I0907 17:52:18.549983 2034 layer_factory.hpp:77] Creating layer loss
I0907 17:52:18.550222 2034 net.cpp:150] Setting up loss
I0907 17:52:18.550245 2034 net.cpp:157] Top shape: (1)
I0907 17:52:18.550249 2034 net.cpp:160] with loss weight 1
I0907 17:52:18.550269 2034 net.cpp:165] Memory required for data: 78018004
I0907 17:52:18.550276 2034 layer_factory.hpp:77] Creating layer accuracy
```

```

I0907 17:52:18.550307 2034 net.cpp:106] Creating Layer accuracy
I0907 17:52:18.550318 2034 net.cpp:454] accuracy <- ip1_ip1_0_split_1
I0907 17:52:18.550325 2034 net.cpp:454] accuracy <- label_data_1_split_1
I0907 17:52:18.550334 2034 net.cpp:411] accuracy -> accuracy
I0907 17:52:18.550349 2034 net.cpp:150] Setting up accuracy
I0907 17:52:18.550362 2034 net.cpp:157] Top shape: (1)
I0907 17:52:18.550367 2034 net.cpp:165] Memory required for data: 78018008
I0907 17:52:18.550372 2034 net.cpp:228] accuracy does not need backward computation.
I0907 17:52:18.550379 2034 net.cpp:226] loss needs backward computation.
I0907 17:52:18.550384 2034 net.cpp:226] ip1_ip1_0_split needs backward computation.
I0907 17:52:18.550390 2034 net.cpp:226] ip1 needs backward computation.
I0907 17:52:18.550395 2034 net.cpp:226] dropip2 needs backward computation.
I0907 17:52:18.550401 2034 net.cpp:226] reluip2 needs backward computation.
I0907 17:52:18.550405 2034 net.cpp:226] ip2 needs backward computation.
I0907 17:52:18.550411 2034 net.cpp:226] pool3 needs backward computation.
I0907 17:52:18.550415 2034 net.cpp:226] relu3 needs backward computation.
I0907 17:52:18.550426 2034 net.cpp:226] conv3 needs backward computation.
I0907 17:52:18.550431 2034 net.cpp:226] pool2 needs backward computation.
I0907 17:52:18.550441 2034 net.cpp:226] relu2 needs backward computation.
I0907 17:52:18.550446 2034 net.cpp:226] conv2 needs backward computation.
I0907 17:52:18.550451 2034 net.cpp:226] relu1 needs backward computation.
I0907 17:52:18.550456 2034 net.cpp:226] pool1 needs backward computation.
I0907 17:52:18.550460 2034 net.cpp:226] conv1 needs backward computation.
I0907 17:52:18.550467 2034 net.cpp:228] label_data_1_split does not need backward computation.
I0907 17:52:18.550475 2034 net.cpp:228] data does not need backward computation.
I0907 17:52:18.550484 2034 net.cpp:270] This network produces output accuracy
I0907 17:52:18.550490 2034 net.cpp:270] This network produces output loss
I0907 17:52:18.550505 2034 net.cpp:283] Network initialization done.
I0907 17:52:18.550590 2034 solver.cpp:60] Solver scaffolding done.
I0907 17:52:18.550948 2034 caffe.cpp:213] Starting Optimization
I0907 17:52:18.550966 2034 solver.cpp:280] Solving
I0907 17:52:18.550971 2034 solver.cpp:281] Learning Rate Policy: step
I0907 17:52:18.551618 2034 solver.cpp:338] Iteration 0, Testing net (#0)
I0907 17:52:20.988834 2034 solver.cpp:406] Test net output #0: accuracy
= 0.4934
I0907 17:52:20.988890 2034 solver.cpp:406] Test net output #1: loss = 0.
693383 (* 1 = 0.693383 loss)
I0907 17:52:21.030238 2034 solver.cpp:229] Iteration 0, loss = 0.692136
I0907 17:52:21.030270 2034 solver.cpp:245] Train net output #0: loss =
0.692136 (* 1 = 0.692136 loss)
I0907 17:52:21.030292 2034 sgd_solver.cpp:106] Iteration 0, lr = 0.01
I0907 17:52:23.047680 2034 solver.cpp:229] Iteration 20, loss = 0.699369
I0907 17:52:23.047734 2034 solver.cpp:245] Train net output #0: loss =
0.699369 (* 1 = 0.699369 loss)
I0907 17:52:23.047744 2034 sgd_solver.cpp:106] Iteration 20, lr = 0.01
I0907 17:52:25.063765 2034 solver.cpp:229] Iteration 40, loss = 0.706948
I0907 17:52:25.063819 2034 solver.cpp:245] Train net output #0: loss =
0.706948 (* 1 = 0.706948 loss)
I0907 17:52:25.063829 2034 sgd_solver.cpp:106] Iteration 40, lr = 0.01
I0907 17:52:27.077561 2034 solver.cpp:229] Iteration 60, loss = 0.673435
I0907 17:52:27.077615 2034 solver.cpp:245] Train net output #0: loss =
0.673435 (* 1 = 0.673435 loss)

```

```

I0907 17:52:27.077625 2034 sgd_solver.cpp:106] Iteration 60, lr = 0.01
I0907 17:52:29.091864 2034 solver.cpp:229] Iteration 80, loss = 0.659732
I0907 17:52:29.091917 2034 solver.cpp:245] Train net output #0: loss =
  0.659732 (* 1 = 0.659732 loss)
I0907 17:52:29.091928 2034 sgd_solver.cpp:106] Iteration 80, lr = 0.01
I0907 17:52:31.105305 2034 solver.cpp:229] Iteration 100, loss = 0.670091
I0907 17:52:31.105362 2034 solver.cpp:245] Train net output #0: loss =
  0.670091 (* 1 = 0.670091 loss)
I0907 17:52:31.105372 2034 sgd_solver.cpp:106] Iteration 100, lr = 0.01
I0907 17:52:33.121857 2034 solver.cpp:229] Iteration 120, loss = 0.673423
I0907 17:52:33.121913 2034 solver.cpp:245] Train net output #0: loss =
  0.673423 (* 1 = 0.673423 loss)
I0907 17:52:33.121923 2034 sgd_solver.cpp:106] Iteration 120, lr = 0.01
I0907 17:52:35.136251 2034 solver.cpp:229] Iteration 140, loss = 0.642096
I0907 17:52:35.136306 2034 solver.cpp:245] Train net output #0: loss =
  0.642096 (* 1 = 0.642096 loss)
I0907 17:52:35.136315 2034 sgd_solver.cpp:106] Iteration 140, lr = 0.01
I0907 17:52:37.153023 2034 solver.cpp:229] Iteration 160, loss = 0.615465
I0907 17:52:37.153086 2034 solver.cpp:245] Train net output #0: loss =
  0.615465 (* 1 = 0.615465 loss)
I0907 17:52:37.153096 2034 sgd_solver.cpp:106] Iteration 160, lr = 0.01
I0907 17:52:39.169787 2034 solver.cpp:229] Iteration 180, loss = 0.579013
I0907 17:52:39.169842 2034 solver.cpp:245] Train net output #0: loss =
  0.579013 (* 1 = 0.579013 loss)
I0907 17:52:39.169852 2034 sgd_solver.cpp:106] Iteration 180, lr = 0.01
I0907 17:52:41.186693 2034 solver.cpp:229] Iteration 200, loss = 0.614775
I0907 17:52:41.186748 2034 solver.cpp:245] Train net output #0: loss =
  0.614775 (* 1 = 0.614775 loss)
I0907 17:52:41.186759 2034 sgd_solver.cpp:106] Iteration 200, lr = 0.01
I0907 17:52:43.203661 2034 solver.cpp:229] Iteration 220, loss = 0.635389
I0907 17:52:43.203716 2034 solver.cpp:245] Train net output #0: loss =
  0.635389 (* 1 = 0.635389 loss)
I0907 17:52:43.203725 2034 sgd_solver.cpp:106] Iteration 220, lr = 0.01
I0907 17:52:45.219781 2034 solver.cpp:229] Iteration 240, loss = 0.610933
I0907 17:52:45.219836 2034 solver.cpp:245] Train net output #0: loss =
  0.610933 (* 1 = 0.610933 loss)
I0907 17:52:45.219846 2034 sgd_solver.cpp:106] Iteration 240, lr = 0.01
I0907 17:52:46.127920 2034 solver.cpp:338] Iteration 250, Testing net (#0)
I0907 17:52:48.642901 2034 solver.cpp:406] Test net output #0: accuracy =
  0.6642
I0907 17:52:48.643168 2034 solver.cpp:406] Test net output #1: loss = 0.
  610271 (* 1 = 0.610271 loss)
I0907 17:52:49.677594 2034 solver.cpp:229] Iteration 260, loss = 0.670394
I0907 17:52:49.677649 2034 solver.cpp:245] Train net output #0: loss =
  0.670394 (* 1 = 0.670394 loss)
I0907 17:52:49.677659 2034 sgd_solver.cpp:106] Iteration 260, lr = 0.01
I0907 17:52:51.695935 2034 solver.cpp:229] Iteration 280, loss = 0.589702
I0907 17:52:51.695987 2034 solver.cpp:245] Train net output #0: loss =
  0.589702 (* 1 = 0.589702 loss)
I0907 17:52:51.696000 2034 sgd_solver.cpp:106] Iteration 280, lr = 0.01
I0907 17:52:53.710996 2034 solver.cpp:229] Iteration 300, loss = 0.517287
I0907 17:52:53.711046 2034 solver.cpp:245] Train net output #0: loss =
  0.517287 (* 1 = 0.517287 loss)
I0907 17:52:53.711056 2034 sgd_solver.cpp:106] Iteration 300, lr = 0.01
I0907 17:52:55.727901 2034 solver.cpp:229] Iteration 320, loss = 0.667897
I0907 17:52:55.727957 2034 solver.cpp:245] Train net output #0: loss =
  0.667897 (* 1 = 0.667897 loss)

```

```
I0907 17:52:55.727967 2034 sgd_solver.cpp:106] Iteration 320, lr = 0.01
I0907 17:52:57.743724 2034 solver.cpp:229] Iteration 340, loss = 0.577245
I0907 17:52:57.743778 2034 solver.cpp:245] Train net output #0: loss =
0.577245 (* 1 = 0.577245 loss)
I0907 17:52:57.743788 2034 sgd_solver.cpp:106] Iteration 340, lr = 0.01
I0907 17:52:59.758558 2034 solver.cpp:229] Iteration 360, loss = 0.577936
I0907 17:52:59.758615 2034 solver.cpp:245] Train net output #0: loss =
0.577936 (* 1 = 0.577936 loss)
I0907 17:52:59.758625 2034 sgd_solver.cpp:106] Iteration 360, lr = 0.01
I0907 17:53:01.775216 2034 solver.cpp:229] Iteration 380, loss = 0.599117
I0907 17:53:01.775274 2034 solver.cpp:245] Train net output #0: loss =
0.599117 (* 1 = 0.599117 loss)
I0907 17:53:01.775286 2034 sgd_solver.cpp:106] Iteration 380, lr = 0.01
I0907 17:53:03.790626 2034 solver.cpp:229] Iteration 400, loss = 0.654552
I0907 17:53:03.790683 2034 solver.cpp:245] Train net output #0: loss =
0.654552 (* 1 = 0.654552 loss)
I0907 17:53:03.790693 2034 sgd_solver.cpp:106] Iteration 400, lr = 0.01
I0907 17:53:05.807268 2034 solver.cpp:229] Iteration 420, loss = 0.56827
I0907 17:53:05.807324 2034 solver.cpp:245] Train net output #0: loss =
0.56827 (* 1 = 0.56827 loss)
I0907 17:53:05.807334 2034 sgd_solver.cpp:106] Iteration 420, lr = 0.01
I0907 17:53:07.822702 2034 solver.cpp:229] Iteration 440, loss = 0.522814
I0907 17:53:07.822758 2034 solver.cpp:245] Train net output #0: loss =
0.522814 (* 1 = 0.522814 loss)
I0907 17:53:07.822768 2034 sgd_solver.cpp:106] Iteration 440, lr = 0.01
I0907 17:53:09.837556 2034 solver.cpp:229] Iteration 460, loss = 0.593677
I0907 17:53:09.837612 2034 solver.cpp:245] Train net output #0: loss =
0.593677 (* 1 = 0.593677 loss)
I0907 17:53:09.837623 2034 sgd_solver.cpp:106] Iteration 460, lr = 0.01
I0907 17:53:11.855217 2034 solver.cpp:229] Iteration 480, loss = 0.547089
I0907 17:53:11.855274 2034 solver.cpp:245] Train net output #0: loss =
0.547089 (* 1 = 0.547089 loss)
I0907 17:53:11.855284 2034 sgd_solver.cpp:106] Iteration 480, lr = 0.01
I0907 17:53:13.771093 2034 solver.cpp:338] Iteration 500, Testing net (#0)
I0907 17:53:16.279290 2034 solver.cpp:406] Test net output #0: accuracy
= 0.6866
I0907 17:53:16.279343 2034 solver.cpp:406] Test net output #1: loss = 0.
606226 (* 1 = 0.606226 loss)
I0907 17:53:16.304071 2034 solver.cpp:229] Iteration 500, loss = 0.564794
I0907 17:53:16.304103 2034 solver.cpp:245] Train net output #0: loss =
0.564794 (* 1 = 0.564794 loss)
I0907 17:53:16.304114 2034 sgd_solver.cpp:106] Iteration 500, lr = 0.001
I0907 17:53:18.319257 2034 solver.cpp:229] Iteration 520, loss = 0.590697
I0907 17:53:18.319314 2034 solver.cpp:245] Train net output #0: loss =
0.590697 (* 1 = 0.590697 loss)
I0907 17:53:18.319324 2034 sgd_solver.cpp:106] Iteration 520, lr = 0.001
I0907 17:53:20.334669 2034 solver.cpp:229] Iteration 540, loss = 0.549887
I0907 17:53:20.334854 2034 solver.cpp:245] Train net output #0: loss =
0.549887 (* 1 = 0.549887 loss)
I0907 17:53:20.334866 2034 sgd_solver.cpp:106] Iteration 540, lr = 0.001
I0907 17:53:22.349474 2034 solver.cpp:229] Iteration 560, loss = 0.540129
I0907 17:53:22.349531 2034 solver.cpp:245] Train net output #0: loss =
0.540129 (* 1 = 0.540129 loss)
I0907 17:53:22.349541 2034 sgd_solver.cpp:106] Iteration 560, lr = 0.001
I0907 17:53:24.365814 2034 solver.cpp:229] Iteration 580, loss = 0.618205
I0907 17:53:24.365869 2034 solver.cpp:245] Train net output #0: loss =
0.618205 (* 1 = 0.618205 loss)
```

```
I0907 17:53:24.365880 2034 sgd_solver.cpp:106] Iteration 580, lr = 0.001
I0907 17:53:26.384941 2034 solver.cpp:229] Iteration 600, loss = 0.5916
I0907 17:53:26.384996 2034 solver.cpp:245] Train net output #0: loss =
0.5916 (* 1 = 0.5916 loss)
I0907 17:53:26.385006 2034 sgd_solver.cpp:106] Iteration 600, lr = 0.001
I0907 17:53:28.401162 2034 solver.cpp:229] Iteration 620, loss = 0.473342
I0907 17:53:28.401217 2034 solver.cpp:245] Train net output #0: loss =
0.473342 (* 1 = 0.473342 loss)
I0907 17:53:28.401227 2034 sgd_solver.cpp:106] Iteration 620, lr = 0.001
I0907 17:53:30.416332 2034 solver.cpp:229] Iteration 640, loss = 0.604995
I0907 17:53:30.416389 2034 solver.cpp:245] Train net output #0: loss =
0.604995 (* 1 = 0.604995 loss)
I0907 17:53:30.416400 2034 sgd_solver.cpp:106] Iteration 640, lr = 0.001
I0907 17:53:32.434772 2034 solver.cpp:229] Iteration 660, loss = 0.511369
I0907 17:53:32.434828 2034 solver.cpp:245] Train net output #0: loss =
0.511369 (* 1 = 0.511369 loss)
I0907 17:53:32.434837 2034 sgd_solver.cpp:106] Iteration 660, lr = 0.001
I0907 17:53:34.448817 2034 solver.cpp:229] Iteration 680, loss = 0.5145
I0907 17:53:34.448881 2034 solver.cpp:245] Train net output #0: loss =
0.5145 (* 1 = 0.5145 loss)
I0907 17:53:34.448900 2034 sgd_solver.cpp:106] Iteration 680, lr = 0.001
I0907 17:53:36.464830 2034 solver.cpp:229] Iteration 700, loss = 0.505878
I0907 17:53:36.464887 2034 solver.cpp:245] Train net output #0: loss =
0.505878 (* 1 = 0.505878 loss)
I0907 17:53:36.464898 2034 sgd_solver.cpp:106] Iteration 700, lr = 0.001
I0907 17:53:38.480952 2034 solver.cpp:229] Iteration 720, loss = 0.545721
I0907 17:53:38.481015 2034 solver.cpp:245] Train net output #0: loss =
0.545721 (* 1 = 0.545721 loss)
I0907 17:53:38.481025 2034 sgd_solver.cpp:106] Iteration 720, lr = 0.001
I0907 17:53:40.497026 2034 solver.cpp:229] Iteration 740, loss = 0.49553
I0907 17:53:40.497089 2034 solver.cpp:245] Train net output #0: loss =
0.49553 (* 1 = 0.49553 loss)
I0907 17:53:40.497100 2034 sgd_solver.cpp:106] Iteration 740, lr = 0.001
I0907 17:53:41.403877 2034 solver.cpp:456] Snapshotting to binary proto file
checkpoints/snapshot_iter_750.caffemodel
I0907 17:53:41.511615 2034 sgd_solver.cpp:273] Snapshotting solver state to
binary proto file checkpoints/snapshot_iter_750.solverstate
I0907 17:53:41.525238 2034 solver.cpp:338] Iteration 750, Testing net (#0)
I0907 17:53:43.960336 2034 solver.cpp:406] Test net output #0: accuracy =
0.7109
I0907 17:53:43.960391 2034 solver.cpp:406] Test net output #1: loss = 0.
553666 (* 1 = 0.553666 loss)
I0907 17:53:43.960400 2034 solver.cpp:323] Optimization Done.
I0907 17:53:43.960404 2034 caffe.cpp:216] Optimization Done.
```

The classification accuracy remains as 71%.

## Results

## Model Evaluation and Validation

We will now deploy our final trained network to perform classification of new images. For all of the training above we used the Caffe command line interface tools. For classification we are going to use Caffe's Python interface. We will first import the Python libraries we require and create some variables specifying the locations of important files.

```
In [25]: import numpy as np
import matplotlib.pyplot as plt
import matplotlib.cm as cm
%matplotlib inline

plt.rcParams['figure.figsize'] = (6.0, 4.0)

# Make sure that caffe is on the python path:
#caffe_root = '../' # this file is expected to be in {caffe_root}/examples
import sys
#sys.path.insert(0, caffe_root + 'python')

import caffe

# Set the right path to your model definition file, pretrained model weights,
# and the image you would like to classify.
MODEL_FILE = '/home/ubuntu/caffe/examples/Dog_Cat_Classification/src/deploy3.prototxt'
PRETRAINED = '/home/ubuntu/caffe/examples/Dog_Cat_Classification/checkpoints/pretrained.caffemodel'
IMAGE_FILE1 = '/home/ubuntu/Data/dog_cat_32/test/cat_236.jpg'
IMAGE_FILE2 = '/home/ubuntu/Data/dog_cat_32/test/dog_4987.jpg'
LABELS_FILE = '/home/ubuntu/Data/dog_cat_32/labels.txt'
labels=open(LABELS_FILE, 'r').readlines()
```

Loading a network is easy. The `caffe.Classifier` method takes care of everything. Note the arguments for configuring input preprocessing: mean subtraction switched on by giving a mean array, input channel swapping takes care of mapping RGB into the reference ImageNet model's BGR order, and raw scaling multiplies the feature scale from the input [0,1] to the ImageNet model's [0,255].

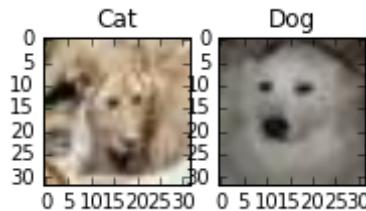
```
In [26]: # First we must import the mean.binaryproto mean image into a numpy array
blob = caffe.proto.caffe_pb2.BlobProto()
data = open('mean.binaryproto', 'rb').read()
blob.ParseFromString(data)
arr = np.array(caffe.io.blobproto_to_array(blob) )
out = arr[0]
```

```
In [27]: # Load our pretrained model
net = caffe.Classifier(MODEL_FILE, PRETRAINED,
                      mean=out,
                      channel_swap=(2,1,0),
                      raw_scale=255,
                      image_dims=(32, 32))
net2 = caffe.Classifier(MODEL_FILE, PRETRAINED,
                      mean=out,
                      channel_swap=(2,1,0),
                      raw_scale=255,
                      image_dims=(32, 32))
```

Let's take a look at our example images with Caffe's image loading helper. We are going to classify 2 different images, one from each category.

```
In [28]: # Load two test images
input_image1 = caffe.io.load_image(IMAGE_FILE1)
input_image2 = caffe.io.load_image(IMAGE_FILE2)
# Display the test images
plt.subplot(1,4,1).imshow(input_image1),plt.title('Cat')
plt.subplot(1,4,2).imshow(input_image2),plt.title('Dog')
```

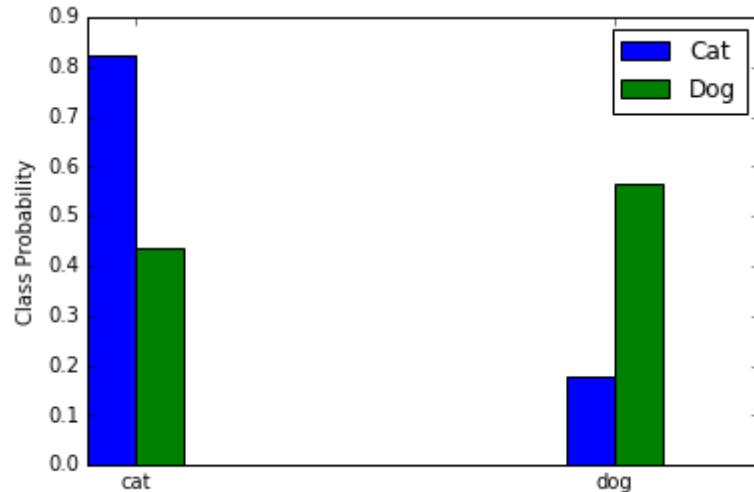
```
Out[28]: (<matplotlib.image.AxesImage at 0x7f33812ff910>,
<matplotlib.text.Text at 0x7f3381314d50>)
```



Time to classify. The default is to actually do 10 predictions, cropping the center and corners of the image as well as their mirrored versions, and average over the predictions. This approach typically leads to better classification accuracy as it is more robust to object translation within in the image.

```
In [29]: prediction1 = net.predict([input_image1])
prediction2 = net2.predict([input_image2])
width=0.1
plt.bar(np.arange(2),prediction1[0],width,color='blue',label='Cat')
plt.bar(np.arange(2)+width,prediction2[0],width,color='green',label='Dog')
plt.xticks(np.arange(2)+width,labels)
plt.ylabel('Class Probability')
plt.legend()
```

Out[29]: <matplotlib.legend.Legend at 0x7f33811d3490>



You can see what class the neural network believes each image is. In the cases above the highest probabilities are given to the correct class for both test images.

## Justification

To validate the performance of our model, a good technique would be to utilize the `sklearn.cross_validation.KFold` module which splits the training set into n-folds and uses n-1 folds for training and 1 fold for testing in a cyclical process that trains and tests every fold. We can compare the n predictions for the folds and see how much they vary.

However, because of the cost associated with training K different models I avoid cross validation. Instead of doing cross validation, we use a random subset of the training data as a hold-out for validation purposes.

A text file containing a list of 20 images of unseen images is provided. By executing the cell below we will classify all of these images with the network we just trained above and calculate the mean accuracy.

```
In [30]: TEST_FILE=open('/home/ubuntu/Data/dog_cat_32/val/test.txt','r')
TEST_IMAGES=TEST_FILE.readlines()
PredictScore=np.zeros((len(TEST_IMAGES),1))
for i in range(len(TEST_IMAGES)):
    IMAGE_FILE='/home/ubuntu/Data/dog_cat_32/val/' + TEST_IMAGES[i].split()[0]
    CATEGORY=TEST_IMAGES[i].split()[1]
    #print TEST_IMAGES[i]
    input_test = caffe.io.load_image(IMAGE_FILE)
    prediction = net.predict([input_test])
    #print prediction[0]
    if prediction[0].argmax()==int(CATEGORY):
        print 'CORRECT -- predicted class for ', str(IMAGE_FILE[62:]),':', prediction[0].argmax(), 'true class:', CATEGORY
    elif prediction[0].argmax()!=int(CATEGORY):
        print 'WRONG -- predicted class ', str(IMAGE_FILE[62:]),':', prediction[0].argmax(), 'true class:', CATEGORY
    PredictScore[i]=int(prediction[0].argmax()==int(CATEGORY))
Accuracy=np.sum(PredictScore)/len(PredictScore)
print 'Prediction accuracy with this image set is', np.sum(PredictScore)/len(PredictScore)
```

```
CORRECT -- predicted class for : 0 true class: 0
CORRECT -- predicted class for : 0 true class: 0
CORRECT -- predicted class for : 0 true class: 0
CORRECT -- predicted class for : 0 true class: 0
CORRECT -- predicted class for : 0 true class: 0
CORRECT -- predicted class for : 0 true class: 0
CORRECT -- predicted class for : 0 true class: 0
WRONG -- predicted class : 1 true class: 0
CORRECT -- predicted class for : 0 true class: 0
CORRECT -- predicted class for : 0 true class: 0
CORRECT -- predicted class for : 1 true class: 1
CORRECT -- predicted class for : 1 true class: 1
CORRECT -- predicted class for : 1 true class: 1
CORRECT -- predicted class for : 1 true class: 1
WRONG -- predicted class : 0 true class: 1
WRONG -- predicted class : 0 true class: 1
CORRECT -- predicted class for : 1 true class: 1
WRONG -- predicted class : 0 true class: 1
WRONG -- predicted class : 0 true class: 1
Prediction accuracy with this image set is 0.75
```

The output above shows our model has correctly classified 75% of our testing set. 75% accuracy is still way below the benchmark (99%) we discussed earlier.

While the solution didn't produce a classifier close to the benchmark given above, it did produce a reasonable classifier which should be very beneficial in classifying Dogs/Cats in images. By evaluating images of the wrong classifications, it appeared to me that those images were either blurry, mix of dogs and cats, an image of big cats (e.g. lions).

It is clear, from output, that this solution model was not sufficient to solve this problem – it seems highly unlikely that even with further refinement this model would perform better than the benchmark. To achieve better accuracy results, more images of different families of cats/dogs required.

## Try your own image

Now we'll grab an image from the web and classify it using the steps above.

- Try setting my\_image\_url to any JPEG image URL.

```
In [31]: # download an image
my_image_url = "http://cdn.playbuzz.com/cdn/0079c830-3406-4c05-a5c1-bc43e8f01479/7dd84d70-768b-492b-88f7-a6c70f2db2e9.jpg" # paste your URL here
!wget -O image.jpg $my_image_url

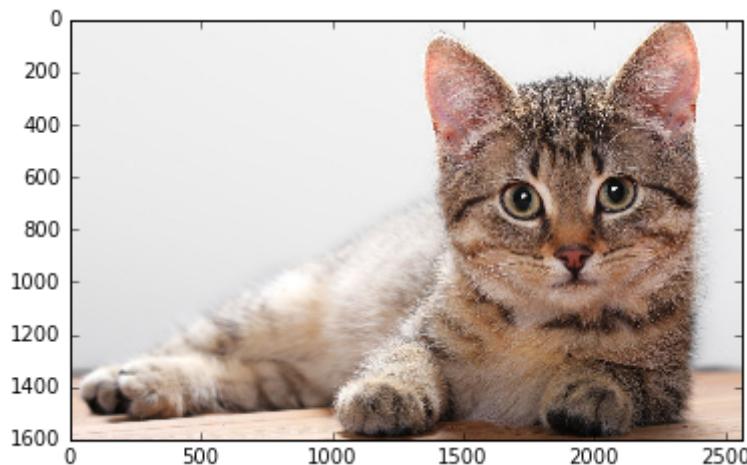
# transform it and copy it into the net
image = caffe.io.load_image('image.jpg')
plt.imshow(image)

--2016-09-07 18:05:05-- http://cdn.playbuzz.com/cdn/0079c830-3406-4c05-a5c1-bc43e8f01479/7dd84d70-768b-492b-88f7-a6c70f2db2e9.jpg
Resolving cdn.playbuzz.com (cdn.playbuzz.com)... 192.229.163.81
Connecting to cdn.playbuzz.com (cdn.playbuzz.com)|192.229.163.81|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 616630 (602K) [image/jpeg]
Saving to: 'image.jpg'

100%[=====] 616,630      --.-K/s   in 0.01s
```

2016-09-07 18:05:05 (50.0 MB/s) - 'image.jpg' saved [616630/616630]

Out[31]: <matplotlib.image.AxesImage at 0x7f3381142990>



```
In [32]: prediction = net.predict([image])
if (prediction[0].argmax() == 1):
    print "This is a DOG!"
else:
    print "This is a CAT!"
```

This is a CAT!

## Conclusion

## Free-Form Visualization

This portion of the project follows the filter visualization example provided with Caffe and the DeCAF visualizations originally developed by Yangqing Jia.

```
In [33]: #View a list of the network layer outputs and their dimensions  
[(k, v.data.shape) for k, v in net.blobs.items()]
```

```
Out[33]: [('data', (1, 3, 32, 32)),  
 ('conv1', (1, 64, 32, 32)),  
 ('pool1', (1, 64, 16, 16)),  
 ('norm1', (1, 64, 16, 16)),  
 ('conv2', (1, 128, 16, 16)),  
 ('pool2', (1, 128, 8, 8)),  
 ('norm2', (1, 128, 8, 8)),  
 ('conv3', (1, 128, 8, 8)),  
 ('pool3', (1, 128, 4, 4)),  
 ('ip2', (1, 500)),  
 ('ip3', (1, 500)),  
 ('ip1', (1, 2)),  
 ('prob', (1, 2))]
```

First you are going to visualize the filters of the first layer.

```
In [34]: # take an array of shape (n, height, width) or (n, height, width, channels)
# and visualize each (height, width) thing in a grid of size approx. sqrt(n) by sqrt(n)
def vis_square(data, padszie=1, padval=0):
    data -= data.min()
    data /= data.max()

    # force the number of filters to be square
    n = int(np.ceil(np.sqrt(data.shape[0])))
    padding = ((0, n ** 2 - data.shape[0]), (0, padszie), (0, padszie)) + ((0, 0),) * (data.ndim - 3)
    data = np.pad(data, padding, mode='constant', constant_values=(padval, padval))

    # tile the filters into an image
    data = data.reshape((n, n) + data.shape[1:]).transpose((0, 2, 1, 3)) + tuple(range(4, data.ndim + 1)))
    data = data.reshape((n * data.shape[1], n * data.shape[3]) + data.shape[4:])

    plt.imshow(data)

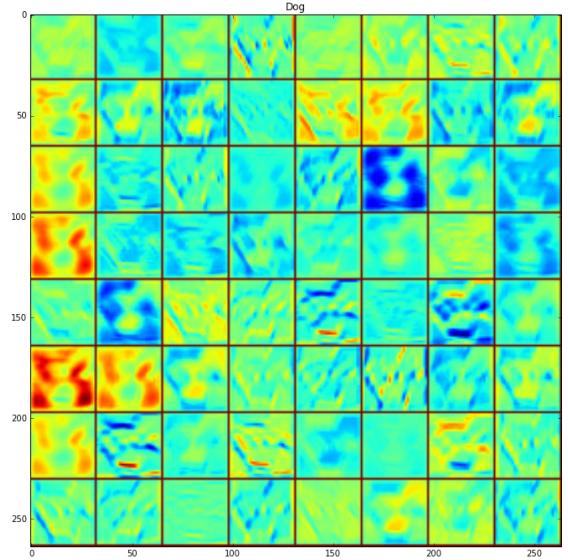
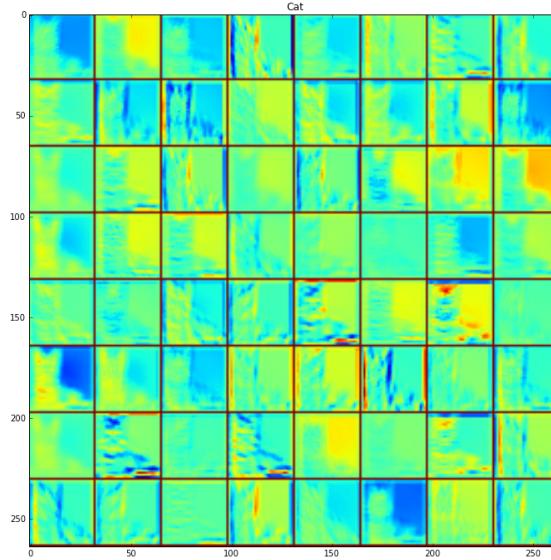
    # the parameters are a list of [weights, biases]

plt.rcParams['figure.figsize'] = (25.0, 20.0)
filters = net.params['conv1'][0].data
vis_square(filters.transpose(0, 2, 3, 1))
```



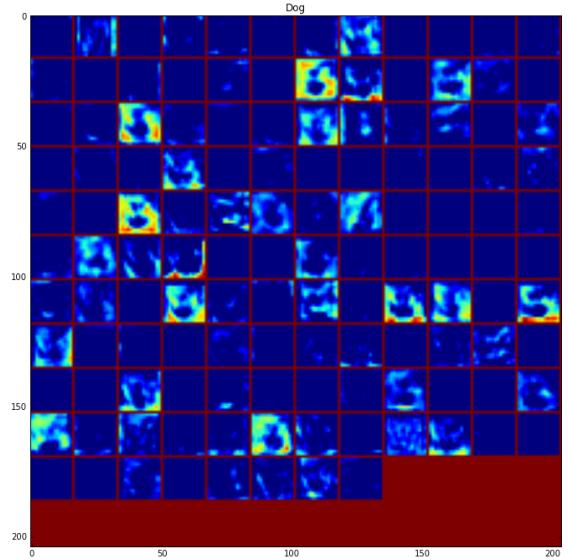
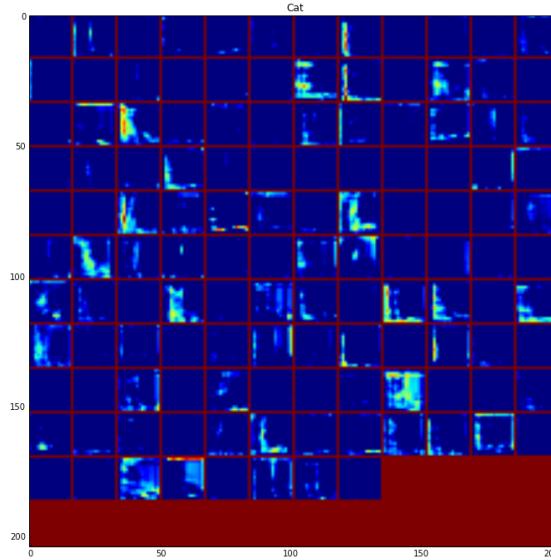
Now you are going to view the feature maps of the two input images after they have been processed by the first convolutional layer. Feel free to modify the `feat` variables so that you can take a closer look as some of the feature maps more closely. Notice the visual similarities and differences between the features maps of both of these images.

```
In [35]: feat = net.blobs['conv1'].data[0,:64]
plt.subplot(1,2,1),plt.title('Cat')
vis_square(feat, padval=1)
net.blobs['conv1'].data.shape
feat2 = net2.blobs['conv1'].data[0, :64]
plt.subplot(1,2,2),plt.title('Dog')
vis_square(feat2, padval=1)
```



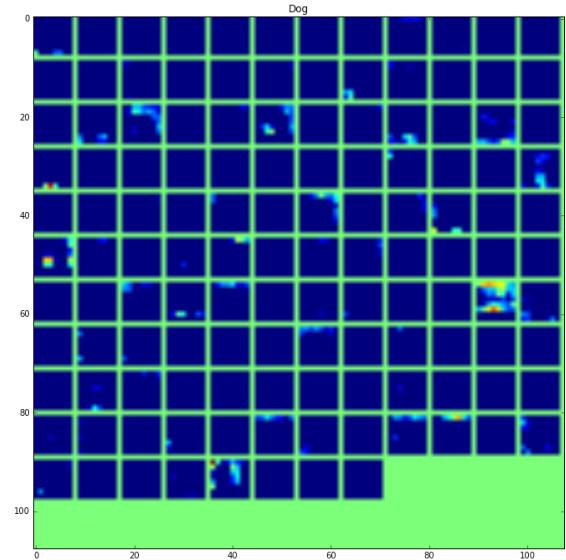
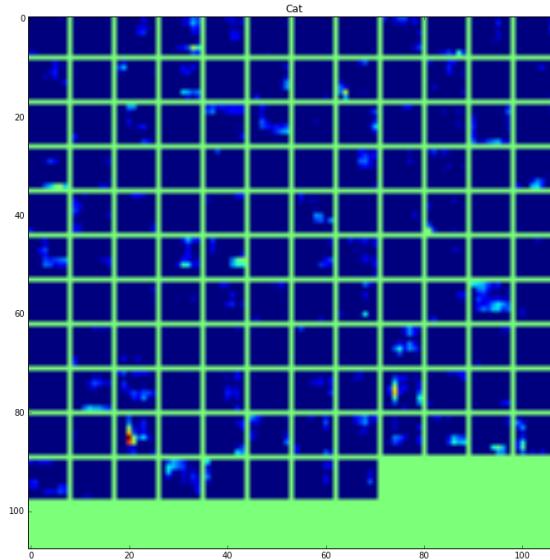
Now view the feature maps from the 2nd convolutional layer.

```
In [36]: feat = net.blobs['conv2'].data[0]
plt.subplot(1,2,1),plt.title('Cat')
vis_square(feat, padval=1)
feat2 = net2.blobs['conv2'].data[0]
plt.subplot(1,2,2),plt.title('Dog')
vis_square(feat2, padval=1)
```

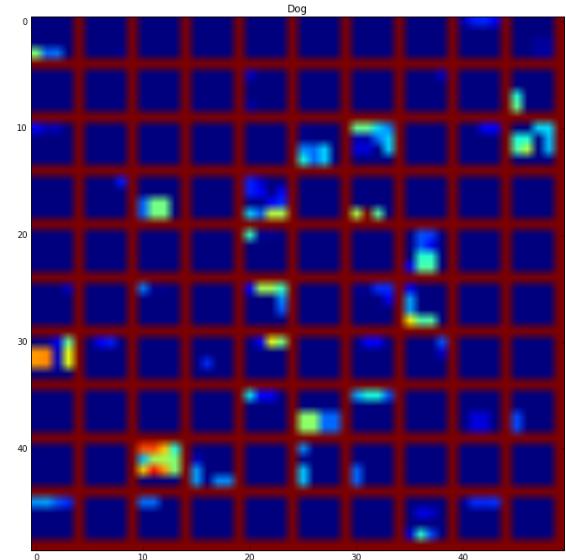
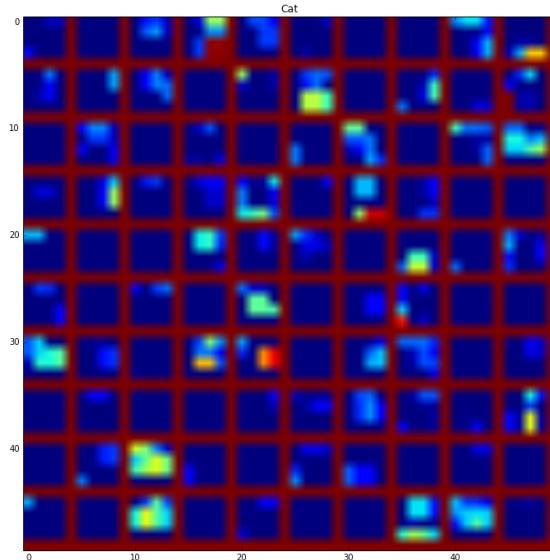


Now view the feature map of the last convolutional layer and then the pooled version.

```
In [37]: feat = net.blobs['conv3'].data[0]
plt.subplot(1,2,1),plt.title('Cat')
vis_square(feat, padval=0.5)
feat2 = net2.blobs['conv3'].data[0]
plt.subplot(1,2,2),plt.title('Dog')
vis_square(feat2, padval=0.5)
```

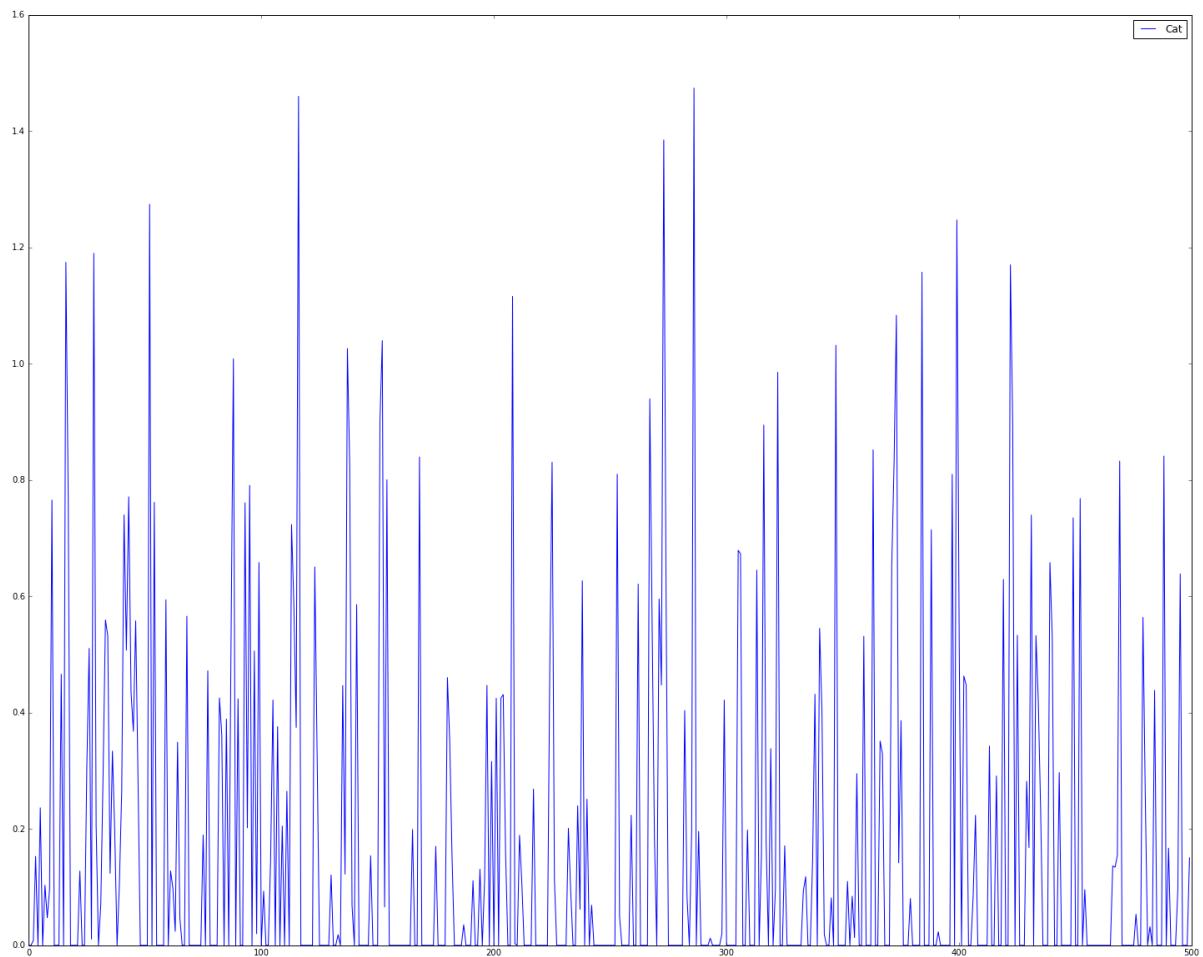


```
In [38]: feat = net.blobs['pool3'].data[0,:100]
plt.subplot(1,2,1),plt.title('Cat')
vis_square(feat, padval=1)
feat2 = net2.blobs['pool3'].data[0,:100]
plt.subplot(1,2,2),plt.title('Dog')
vis_square(feat2, padval=1)
```

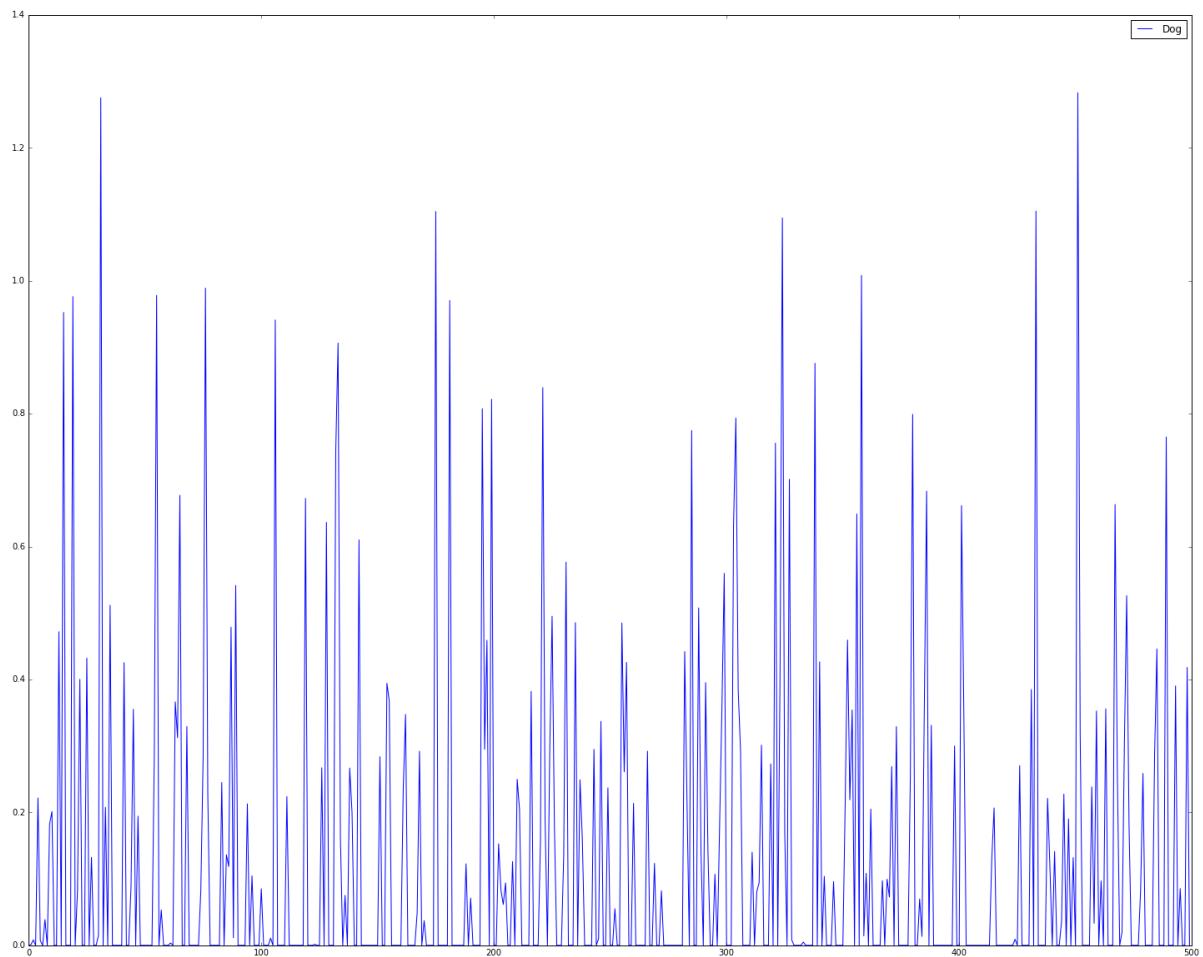


Now view the neuron activations for the fully-connected layer ip2. You will notice that the neurons being activated by the two input images are very different. This is good as it means the network is effectively differentiating the two images at the higher layers in the network.

```
In [39]: feat = net.blobs['ip2'].data[0]
plt.plot(feat.flat,label='Cat')
plt.legend()
plt.show()
#####
#####
# Plot ip2 for the input image of the Dog image. Compare the Differences
feat2 = net2.blobs['ip2'].data[0]
plt.plot(feat2.flat, label='Dog')
plt.legend()
```



Out[39]: <matplotlib.legend.Legend at 0x7f3e888c9ad0>



## Reflection

I got interested in object recognition in images when I first watched a PyData video on YouTube. PyData is a gathering of users and developers of data analysis tools in Python, and there is PyData channel on YouTube with lots of videos on different Machine Learning subjects. This particular video that got me interested was from PyData Berlin 2016 and it was about Brand recognition in real-life photos where the presenter talked about the project he worked on to detect brands in photos posted on Instagram. The video link is [here].(<https://www.youtube.com/watch?v=0Ho0O1tvcU4>) (<https://www.youtube.com/watch?v=0Ho0O1tvcU4>)

I got the source code for the project but unfortunately the python notebook did not run on my Windows laptop due to missing libraries. Also I did not have access to the dataset used in this project.

I started looking for brand logos dataset and came across [FlickrLogos-32\\_dataset\\_v2](http://www.multimedia-computing.de/flickrlogos/) (<http://www.multimedia-computing.de/flickrlogos/>). I downloaded the file and started playing with the dataset.

Next, I started studying Caffe. I soon realized that Caffe isn't like scikitlearn: it hasn't been designed so it easily predict using `model.fit(X, y)` or `model.predict(Z)` and get a useful result.

I should warn you that Caffe has a steep learning curve, especially if you expect to use it wisely and do some changes and customization. In that level of detail, its documentation is lacking as well (this is of course my personal opinion).

Installation of Caffe is also a challenge since it is required to download the source code and compile the code on a Linux machine. Caffe has several dependencies like CUDA for GPU mode that needs to be installed if you are planning to use GPU. I also learned that there are a lot of vector and matrix operations in deep learning, so it's intuitive that deep learning should run several times faster on a GPU than a CPU just like a game runs faster on a GPU compared to a CPU.

In short, it is a big hassle to install Caffe, and my recommendation is don't even try it to install it on your Windows laptop.

Instead of installing Caffe from scratch, I got a AWS instance with GPU support and started with an AMI that has already installed with all the required software.

I started with Caffe tutorials and used my custom logo dataset. On my first try my jupyter server crashed. I started researching and reading more articles to make this work but Unfortunately it was not a successful experience.

I then decided to switch the project from brand recognition in photos to Dogs/Cats image classifier which I also found it interesting. This was also a competition at Kaggle. I tried to learn as much as I could with this project and then go back to Brand Recognition problem. Hopefully someday you will see this brand recognition project on my Github repo.

The field is moving very fast and new and interesting architectures and techniques are been discussed and used all the time. Because of the success of CNNs in the field of image classification, in many different areas where convolutional neural networks are applied to solve ultra-complex problems, often the input data first is translated to image data.

Understanding ConvNets and learning to use them for the first time can sometimes be an intimidating experience but I found Deep Learning and especially Convolutional Neural Networks very interesting, and there are lots of applications like Image recognition, Video analysis, Natural language processing, Drug discovery, and Playing Go!

## Improvement

My Goal in this project was to create Convolutional Neural Network from scratch, however, lots of researchers and engineers have made Caffe models for different tasks with all kinds of architectures and data. These models are learned and applied for problems ranging from simple regression, to large-scale visual classification, to Siamese networks for image similarity, to speech and robotics applications.

To improve the classification accuracy, we might increase network size or use one of the Caffe's predefined models from [Caffe Model Zoo](http://caffe.berkeleyvision.org/model_zoo.html) ([http://caffe.berkeleyvision.org/model\\_zoo.html](http://caffe.berkeleyvision.org/model_zoo.html)), change activation functions, modify training algorithm, and do data augmentation.

Another possibility is using augmented image generators, such as in keras with the `ImageDataGenerator` class, in order to flip, rotate, and transform your old images to produce new images that can add to your training set.

## References

1. [Caffe - Deep learning framework by the BVLC](http://caffe.berkeleyvision.org/) (<http://caffe.berkeleyvision.org/>)
2. [Wikipedia - Convolutional Neural Network](https://en.wikipedia.org/wiki/Convolutional_neural_network)  
([https://en.wikipedia.org/wiki/Convolutional\\_neural\\_network](https://en.wikipedia.org/wiki/Convolutional_neural_network))
3. [An Intuitive Explanation of Convolutional Neural Networks](https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/) (<https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>)
4. [Deep Learning in a Nutshell: Core Concepts](https://devblogs.nvidia.com/parallelforall/deep-learning-nutshell-core-concepts/) (<https://devblogs.nvidia.com/parallelforall/deep-learning-nutshell-core-concepts/>)
5. [Start deep learning with Jupyter notebooks in the cloud](http://efavdb.com/deep-learning-with-jupyter-on-aws/) (<http://efavdb.com/deep-learning-with-jupyter-on-aws/>)
6. [Udacity Deep Learning - Take machine learning to the next level](https://www.udacity.com/course/deep-learning--ud730)  
(<https://www.udacity.com/course/deep-learning--ud730>)
7. [Installation Guide - Caffe and Anaconda on AWS EC2](https://github.com/adilmoujahid/deeplearning-cats-dogs-tutorial/blob/master/aws-ec2-setup.md) (<https://github.com/adilmoujahid/deeplearning-cats-dogs-tutorial/blob/master/aws-ec2-setup.md>)
8. [CS231n Convolutional Neural Networks for Visual Recognition](http://cs231n.github.io/neural-networks-3/) (<http://cs231n.github.io/neural-networks-3/>)

In [ ]: